Implementation Guide

Connected Mobility Solution on AWS



Copyright © 2024 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Connected Mobility Solution on AWS: Implementation Guide

Copyright © 2024 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

| Solution overview | |
|--|----|
| Features and benefits | 3 |
| Use cases | 6 |
| Concepts and definitions | 7 |
| Architecture overview | 10 |
| Automotive Cloud Developer Portal deployment | 10 |
| Networking scenarios | 11 |
| Deploying CMS on AWS modules with Backstage | 12 |
| CMS on AWS modules and services | 13 |
| Amazon Virtual Private Cloud (Amazon VPC) | 13 |
| Auth Setup | 13 |
| Config | 14 |
| Automotive Cloud Developer Portal (ACDP) and Backstage | 14 |
| Auth | 14 |
| AWS IoT Core and MQTT | 14 |
| Vehicle Provisioning | 15 |
| Connect and Store | 15 |
| FleetWise Connector | 15 |
| API | 15 |
| Alerts | 16 |
| EV Battery Health | 16 |
| Vehicle Simulator | 16 |
| Predictive Maintenance | 16 |
| Fleet Management UI (Preview) | 17 |
| AWS Well-Architected design considerations | 17 |
| Operational excellence | 17 |
| Security | 17 |
| Reliability | 18 |
| Performance efficiency | 18 |
| Cost optimization | 19 |
| Sustainability | 19 |
| Architecture details | 20 |
| Config | 20 |
| Virtual Private Cloud (VPC) | 21 |

| | Auth Setup | 23 |
|----|---|----|
| | Automotive Cloud Developer Portal | 25 |
| | Supported networking scenarios for ACDP | 26 |
| | Multi Account/Multi Region Deployment | 26 |
| | Backstage module | 28 |
| | Auth module | 31 |
| | Vehicle Provisioning module | 33 |
| | Connect and Store module | 35 |
| | API module | 37 |
| | Alerts module | 38 |
| | EV Battery Health module | 40 |
| | Vehicle Simulator module | 43 |
| | FleetWise Connector module | 45 |
| | Predictive Maintenance module | 45 |
| | Fleet Management UI module | 51 |
| | AWS services in this solution | 53 |
| Ρl | an your deployment | 58 |
| | Supported AWS Regions | 58 |
| | Cost | 58 |
| | CMS on AWS static cost tables | 58 |
| | Example costs | 62 |
| | Security | 69 |
| | Authentication and authorization | 70 |
| | Amazon API Gateway | 70 |
| | Amazon Bedrock | 70 |
| | Content Security Policy (CSP) | 71 |
| | Amazon CloudFront | 71 |
| | Amazon CloudWatch | 71 |
| | Customer managed AWS KMS keys | 71 |
| | Amazon DynamoDB | 72 |
| | Elastic Load Balancing | 72 |
| | PII data | 72 |
| | Amazon SageMaker Al | 72 |
| | AWS WAF | 72 |
| | Other AWS services | 73 |
| | Quotas | 79 |

| Quotas for AWS Services in this solution | 79 |
|---|-----|
| AWS CloudFormation quotas | 79 |
| Limitations | 79 |
| Deploy the solution | 80 |
| Deployment process overview | 80 |
| CMS module deployment order | 81 |
| AWS CloudFormation templates | 82 |
| Monitor the solution with Service Catalog AppRegistry | 99 |
| myApplications dashboard | 99 |
| Activate CloudWatch Application Insights | 100 |
| Confirm cost tags associated with the solution | 101 |
| Activate cost allocation tags associated with the solution | 102 |
| AWS Cost Explorer | 103 |
| Update the solution | 104 |
| Troubleshooting | 105 |
| Problem: Lambda Runtime not supported | 105 |
| Resolution | 105 |
| Problem: Multiple ProvisionedVehicles active certificates | 105 |
| Resolution | 105 |
| Problem: Redirect mismatch error when using Cognito Instance deployed by Auth-Setup | |
| module | 106 |
| Resolution | 106 |
| Contact AWS Support | 106 |
| Create case | 106 |
| How can we help? | 107 |
| Additional information | 107 |
| Help us resolve your case faster | 107 |
| Solve now or contact us | 107 |
| Uninstall the solution | 108 |
| Capture the deployment UUID | 108 |
| Delete CMS on AWS modules in order | 108 |
| Delete the Backstage ACM certificate (optional) | 109 |
| Manually clean up resources | 109 |
| Deleting the Amazon S3 buckets | 110 |
| Deleting the Amazon DynamoDB tables | 111 |
| Deleting the Amazon CloudWatch logs | 111 |

| Deleting the AWS KMS customer managed keys | 112 |
|---|-----|
| Deleting the Amazon Cognito user pools | 112 |
| Deleting the Amazon Relational Database Service snapshots | 113 |
| Developer guide | 114 |
| Source code | 114 |
| Integrating custom modules | 114 |
| Global system integrator (GSI) and independent software (ISV) integration | 114 |
| Using CMS as a foundation | 114 |
| Integrating with existing systems | 115 |
| White-labeling CMS for end customers | 115 |
| Extending functionality with custom modules | 116 |
| Managing deployment and operations | 116 |
| Collaborating with AWS for customizations | 116 |
| Reference | 118 |
| Anonymized data collection | 118 |
| Contributors | 118 |
| Revisions | 119 |
| Nations | 120 |

Accelerate development and deployment of connected vehicle assets

Publication date: October 2023 (last update: April 2025)

Amazon Web Services (AWS) automotive customers have asked for ways to manage fleets with increased efficiency and reduced vehicle downtime through preventative maintenance, location tracking, improved fleet safety and security, and new software driven vehicle experiences.

The Connected Mobility Solution (CMS) on AWS addresses these needs, and provides various capabilities for vehicles and customers to interact with the AWS Cloud. This solution allows you to take advantage of provided features as modules, deploy and manage these modules from a centralized platform, and implement and integrate your own custom modules into the solution.

Provided features allow you to:

- Communicate between vehicles and the AWS Cloud.
- Manage and orchestrate CMS on AWS deployments from a centralized developer platform.
- Securely authenticate and authorize users and services.
- Onboard vehicles into AWS IoT Core, creating profiles for provisioned vehicles.
- Capture, store, and query telemetry data emitted from provisioned vehicles.
- Create alerts for this data, and subscribe to notifications based on data thresholds.
- Visualize vehicle telemetry data through an Amazon Managed Grafana dashboard.
- Simulate connected vehicle data emission.
- Add additional modules for your unique use cases.

CMS on AWS implements an opinionated deployment mechanism for managing a connected vehicle platform. Original equipment manufacturers (OEMs), tier one suppliers, and fleet operators can manage and configure all or a subset of these capabilities for an end-to-end connected vehicle solution, as well as integrate custom module implementations.

Vehicle telemetry data, such as speed, oil temperature, tire pressure, and geolocation generated from car sensors can provide near real-time data ingestion for analytics and machine learning (ML) use cases. CMS on AWS makes this data available for external consumption to third parties. You

1

can leverage charging mechanisms and role-based authorization for usage-based insurance scores, connected accident advisor, and package delivery services.

CMS on AWS provides a modular catalog approach that lets you independently enable workloads specific to your use cases. You can deploy and configure workloads independently or as a whole system to begin securely and economically deriving insights from your connected vehicle data.

This implementation guide provides an overview of CMS on AWS, its reference architecture and components, considerations for planning the deployment, and configuration steps for deploying CMS on AWS to the AWS Cloud.

The intended audience for using this solution's features and capabilities in their environment includes solutions architects, business decision makers, DevOps engineers, data scientists, and cloud professionals.

Use this navigation table to quickly find answers to these questions:

| If you want to | Read |
|---|------------------------------|
| Know the cost for running this solution. | Cost |
| The estimated cost for running this solution in the us-east-1 Region is USD \$352.80 per month. | |
| Understand the security considerations for this solution. | Security |
| Know how to plan for quotas for this solution. | Quotas |
| Know which AWS Regions are supported for this solution. | Supported AWS Regions |
| View or download the AWS CloudForm ation template included in this solution to automatically deploy the infrastructure resources (the "stack") for this solution. | AWS CloudFormation templates |

| If you want to | Read |
|---|-------------------|
| Access the source code and optionally use the AWS Cloud Development Kit (AWS CDK) to deploy the solution. | GitHub repository |

Features and benefits

The solution provides the following features:

Modularity

CMS on AWS simplifies maintaining and implementing the solution by utilizing a modular design. Modularity provides two significant benefits. First, it separates concerns and outlines clear boundaries between functional sets, enabling the solution to flexibly meets connected vehicle system demands. Second, it allows users to interchange CMS on AWS modules with bespoke implementations that meet the boundary definitions defined by CMS on AWS standards, as well as implementing additional modules into the solution. These custom modules can then be deployed from the same portal (ACDP / Backstage) as CMS on AWS provided modules.

Automotive Cloud Developer Portal

Deploying a complex modular system demands a well-defined and streamlined method for selecting, configuring, and deploying each module. This component uses Backstage, AWS
CodeBuild, and AWS CloudFormation to deploy and manage CMS on AWS modules. It uses well-defined templates to allow for implementations of custom modules to be integrated into the solution and deployed from the same portal.

Multi-Account Multi-Region Deployments

The Automotive Cloud Developer Portal provides an optional feature that allows user to deploy modules through Backstage into other AWS accounts and regions. Users can select target accounts and regions at deploy time. This is supported by AWS Control Tower, and required a more advanced deployment structure to support cross-account authorization.

Configuration

CMS on AWS provides a mechanism to associate deployments of CMS on AWS modules with a unique ID. Modules sharing the same unique ID also share infrastructure such as VPC and identity

Features and benefits 3

provider (IdP). Other CMS on AWS modules deployed with the unique ID can use the configuration module to lookup VPC and IdP configurations. The configuration module, which reports metrics based on Amazon Simple Storage Service (Amazon S3) and Amazon Timestream usage handles anonymized metrics collection.

OAuth 2.0 Authentication

CMS on AWS provides a default deployment of authentication infrastructure through Amazon
Cognito
or the ability to integrate with customer's own IdP services. The IdP chosen can be used for both CMS on AWS users and services to authenticate by retrieving access tokens and JSON Web Tokens (JWTs). CMS on AWS then provides the means to validate the integrity of access tokens with the chosen IdP. The chosen IdP is also used to authenticate (e.g. sign in) users to the Backstage portal.

Role-Based Access Control Authorization

The ACDP / Backstage deployment within CMS on AWS includes authorization by default that protects access to all parts of the Backstage system, including for authenticated users. Users can be granted permissions, defined by Backstage and plugin creators, to access different parts of the Backstage frontend and backend APIs. Permissions are grouped into roles, which can be assigned to users and groups, all via a user friendly UI. Although advanced user and group creation and management is not yet supported, all authenticated users who successfully sign-in to the Backstage portal will have a user created for them which can be given authorization permissions.

Network Security

CMS on AWS provides a default VPC or the ability to bring your own VPC to secure the AWS resources deployed in your account. The default VPC provides public, private and isolated subnets in two Availability Zones (AZs). This enables a highly available and secure cloud network. Modules then use these subnets to manage security accordingly.

Vehicle Provisioning

CMS on AWS registers vehicles as AWS IoT Core things to help you securely monitor vehicles, their certificates, and their policies. Vehicle provisioning begins by registering with a claim certificate that the solution generates during deployment. The claim certificate includes a provisioning template that configures AWS IoT Core thing, certificate, and policy creation. After the certificate is registered, the solution provides provisioned vehicles an individual certificate and public/private key pair, which you can use to connect repeatedly in the future. This is the process defined by fleet provisioning.

Features and benefits

Storage

This solution provides a simple storage mechanism for simulated and provisioned vehicle data with Amazon S3. The solution ingests data from pre-defined Message Queuing Telemetry Transport (MQTT) topics, and stores it in both JSON and Parquet data formats. The storage is integrated with an alerts mechanism which utilizes Amazon Simple Notification Service (Amazon SNS).

API

CMS on AWS provides the ability to query vehicle data stored within the solution for use in other CMS and customer-built modules. CMS on AWS uses an <u>AWS AppSync</u>https://graphql.org/learn/[GraphQL] application programming interface (API) that builds and runs <u>Amazon Athena</u> queries to provide near real-time data directly from the CMS on AWS data lake (an Amazon S3 bucket).

Alerts

CMS on AWS provides the ability to send alerts based on customizable user subscriptions through AWS AppSync GraphQL API operations and Amazon SNS. These alerts can be triggered from other CMS on AWS modules, and sent to the user through configurable notification settings.

Electric Vehicle Battery Health

CMS on AWS provides the ability to visualize Electric Vehicle (EV) battery telemetry data and configure alerts based on data thresholds. This is done through an Amazon Managed Grafana. EV battery telemetry data is obtained by running Amazon Athena queries through the CMS API module.

Simulation

For both developers and customers, it is important to have a method for generating simulated vehicle telemetry data. This allows for testing the solution in real time, while also generating datasets. CMS on AWS provides a method for simulating up to 10 vehicles at once, with configurations for how the data is generated (interval, amount, and duration) and a customizable schema to define the payload that is generated.

Metric and Cost Monitoring

This solution includes a <u>Service Catalog AppRegistry</u> resource to register the solution's CloudFormation template and its underlying resources as an application in both Service Catalog AppRegistry and Application Manager. With this integration, you can use myApplications, which is

Features and benefits 5

an extension of the AWS Management Console home. You can view this solution in myApplications to help you to centrally manage the solution's resources and enable application search, reporting, and management actions. You can view metrics such as the overall cost of your deployment, and set up additional monitoring with services like <u>AWS Cost Explorer</u>, <u>Amazon CloudWatch</u>, and <u>AWS Security Hub</u>.

Predictive Maintenance

CMS on AWS provides an MLOps infrastructure that offers the capability to train, retrain, evaluate, and deploy ML models for fault prediction use cases. Additionally, you can run real-time inference and batch inference on the deployed ML model. CMS on AWS also provides a generative artificial intelligence (AI) agent that can answer natural language queries regarding the fleet. This agent uses textual data stored in an S3 bucket and queries the backend for real-time vehicle maintenance status.

Fleet management UI

CMS on AWS provides a user interface that offers the capability to manage vehicles and fleets in AWS IoT FleetWise. Additionally, you can build out widgets and insert them into the UI to monitor your fleet and vehicle via captured telemetry.

Use cases

Deployment and management

Fleet owners and deployment managers can select, configure, and deploy CMS on AWS modules through the centralized Automotive Cloud Developer Portal (ACDP). From here, deployments can also be monitored and tore down. Custom module implementations can be integrated into Backstage, the frontend component of the ACDP, and managed in the same manner. A centralized platform allows for managing and monitoring your deployment easily and effectively, as well as separating ownership and permissions between members of your organization. The ACDP Backstage deployment also supported cross-account and cross-region deployments, given the necessary configuration.

Vehicle connection

Vehicle manufacturers, fleet owners, and vehicle owners can securely register and receive unique credentials by using the <u>fleet provisioning by claim</u> process. You can use these credentials for future connections of the vehicle and securely manage them with AWS IoT Core. This facilitates

Use cases 6

disabling and replacing vehicle credentials at any time. Connected vehicles can then emit data to AWS IOT Core MQTT topics for the solution to ingest and store.

Query data

Vehicle manufacturers, fleet owners, and vehicle owners can leverage flexible queries using GraphQL for data analysis and visualization, ML applications, and stateful representation of individual vehicles. Data is available in near real-time as it is emitted from connected vehicles. Future modules and customer applications can use this feature to realize value from the collected data.

Monitor vehicles

Vehicle manufacturers, owners, and fleet managers can monitor vehicle data through visualizations and alerts available through a configured Amazon Managed Grafana dashboard. Alerts can query vehicle data periodically and send notifications to users when alert thresholds are breached. Visualization is updated in near real-time and you can configure alerts to be evaluated at a desired interval. Furthermore, users can monitor data with Amazon SNS by using the CMS Alerts module, which you can configure with custom alert conditions and functionality from pre-built or custom CMS on AWS modules.

Vehicle simulation

Data and entity simulation is a valuable feature for the development and use of connectivity solutions. The CMS Vehicle Simulator module allows simulating vehicle provisioning and registration (onboarding), data generation, and data ingestion by publishing to MQTT topics. Connected vehicle engineers can use the CMS Vehicle Simulator module to test the solution in a variety of usage scenarios, as well as generate datasets. Fleet managers can use simulations to showcase functionality and ensure the effectiveness of their solution's configuration.

Development monitoring

After you deploy CMS on AWS, you can monitor the modules from the Backstage portal. Deployment managers can monitor their CodeBuild project results and executions, as well as the total cost of their deployment. You can also use the myApplications dashboard and configure it for services such as Cost Explorer, CloudWatch, and Security Hub.

Concepts and definitions

This section describes key concepts and defines terminology specific to this solution:

Concepts and definitions

alert

Alerts refer to any notification sent to a user because of rules set up on stored telemetry data, or triggers from CMS on AWS modules. Alert functionality is currently provided through both the CMS Alerts module (Amazon SNS alerts), and CMS EV Battery Health module (Amazon Managed Grafana alerts).

Backstage

Backstage is an open-source project used by CMS on AWS. It provides an authenticated portal for managing deployments of CMS on AWS modules with the <u>ACDP</u>. It also allows for integration with bespoke customer modules for a centralized deployment mechanism.

fleet provisioning by claim

<u>Fleet provisioning by claim</u> is a provisioning method that is used to authenticate and register vehicles the first time they connect to AWS IoT Core.

Grafana

Grafana is a popular open source visualization and dashboarding platform that allows users to query, visualize, alert on, and understand data. Amazon Managed Grafana is a fully managed service for Grafana. CMS on AWS uses Amazon Managed Grafana to provide a visualization and alert dashboard for vehicle telemetry data.

identity provider

CMS on AWS provides the ability to integrate the authentication system with any third-party identity provider which is OAuth 2.0 compliant. This identity provider should provide some concept of users, a resource server (audience), domain (issuer), and clients for users and services.

Message Queuing Telemetry Transport (MQTT)

<u>MQTT</u> is an OASIS standard messaging protocol used by AWS IoT Core. It is designed as an extremely lightweight publish-subscribe messaging transport.

provisioning

Provisioning refers to the process of registering a vehicle in AWS IoT Core by using a recognized claim certificate. Registered vehicles receive a unique certificate and credentials to connect again in the future.

Concepts and definitions 8

simulation

Simulation refers to the capability of the CMS Vehicle Simulator module to simulate vehicles emitting telemetry data. A single simulation can produce randomized data for multiple simulated vehicles at a time, and emit data at regular intervals for a specified duration.

tier one supplier

A tier one supplier manufactures and supplies original equipment manufacturer (OEM) companies with components. They supply parts and devices that an OEM needs in order to complete a product.

Vehicle Signal Specification

Vehicle Signal Specification (VSS) is a common data schema for standardizing vehicle data formats.

VPC / CMS on AWS VPC

A virtual private cloud (VPC) is an AWS service. Within the context of our solution, VPC is often referring to the specific VPC instance chosen to be deployed with the solution. This is either the default CMS on AWS VPC, a preconfigured VPC provided by CMS on AWS, or a customer provided VPC. Any of these can then be attached to this solution's configuration. In all cases, you must create a VPC prior to deploying CMS on AWS, which is provided to subsequent module deployments for network security. All modules that can use a VPC utilize the same configured VPC.

white-label

Rebranding a product after acquisition.



Note

For a general reference of AWS terms, see the AWS Glossary.

Concepts and definitions

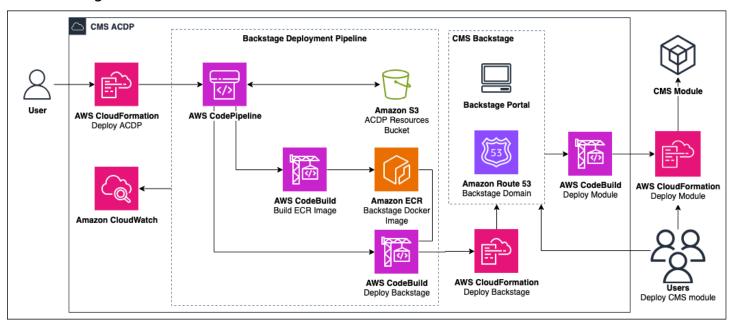
Architecture overview

This section provides a high-level description of all CMS on AWS provided modules. It also includes architecture diagrams for the deployment of the Automotive Cloud Developer Portal and the deployment of subsequent CMS on AWS modules.

Outside of deployment resources and necessary configuration modules, each installation of CMS on AWS is unique based on which provided modules you deploy, and which bespoke module implementations you integrate. The deployment architecture described in this section is consistent across deployments of CMS on AWS.

Automotive Cloud Developer Portal deployment

Deploying this solution with the default parameters deploys the following components of ACDP and Backstage.



Note

CloudFormation resources are created from AWS Cloud Development Kit (AWS CDK) constructs.

The high-level process flow for the solution components deployed with AWS CDK is as follows (from left to right in the diagram):

- Automotive Cloud Developer Portal To aid in orchestrating deployments of CMS on AWS
 modules, you first deploy the ACDP. This deployment is handled with <u>CloudFormation</u> templates
 created by AWS CDK. The ACDP creates a Backstage deployment pipeline. The pipeline then
 deploys the <u>Backstage module</u>. Refer to the <u>Networking</u> file on the GitHub repository for
 information about ways the ACDP can be configured for public internet access, private network
 only access, and access with <u>AWS Transit Gateway</u>.
- 2. **Multi Account/Multi Region Deployment** When deploying the ACDP cloudformation template, ACDP can be configured to enable Multi Account/Multi Region Deployment by providing values for associated CloudFormation parameters. Refer to the <u>ACDP Instructions</u> in the GitHub repository. Upon successful deployment, additional configuration is required in the Organizations management account. Refer to the <u>Guidance</u> in the GitHub repository.
- 3. **Backstage Deployment Pipeline** The ACDP deploys a CI/CD pipeline through <u>AWS</u>

 <u>CodePipeline</u> that configures all the steps necessary to deploy the Backstage module. This is accomplished with <u>AWS CodeBuild</u> pipeline projects, which use build specification files to define their actions. <u>Amazon Elastic Container Registry</u> (Amazon ECR) is used to store the Backstage <u>Docker</u> image. For details on the structure of the pipeline and each build step, see the <u>Backstage</u> module.
- 4. **Backstage** The Backstage module is the presentation layer for the ACDP. An <u>Elastic Load Balancing</u> (ELB) Application Load Balancer connects with <u>Amazon Route 53</u> and an <u>Amazon Elastic Container Service</u> (Amazon ECS) cluster group setup with <u>AWS Fargate</u> tasks. The Backstage module allows deploying CMS on AWS modules through a graphical user interface. For more information, see <u>Backstage module</u>.
- 5. **Deploying CMS on AWS modules with Backstage** When the ACDP is configured, there are two ways to deploy CMS on AWS modules (see Deploy the solution for details):
 - Backstage By using ACDP and Backstage, you can deploy CMS on AWS modules using module templates configured for Backstage, powered by AWS CodeBuild.
 - AWS CDK CLI Without Backstage, you can individually deploy the CMS on AWS modules directly from the repository by utilizing the make deploy target.

Networking scenarios

The ACDP's networking can be configured in the following ways:

- 1. Fully public with Route53 DNS and AWS Certificate Manager (ACM) auto-generated certificate.
- 2. Fully public with external DNS and a customer uploaded TLS certificate to ACM.

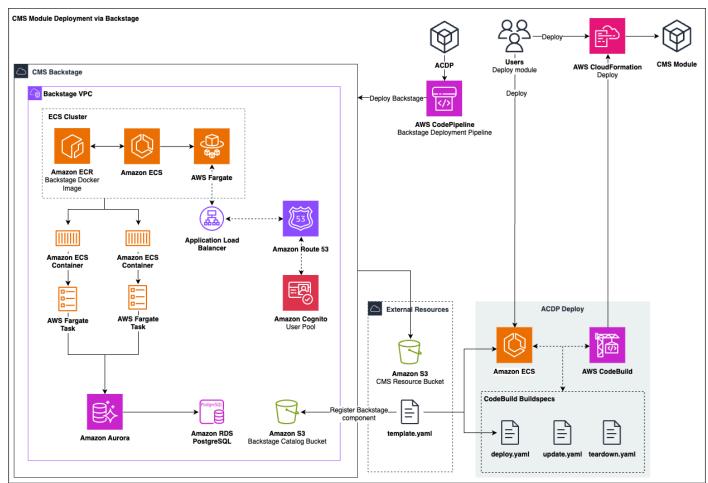
Networking scenarios 11

- 3. Public DNS with internal only access through a bastion or VPC peering.
- 4. Private DNS with Route53 private hosted zone, internal only access through a bastion or VPC peering.
- 5. Private VPC linked to a public VPC through a transit gateway. Access controlled by the public VPC.

Refer to the <u>Networking</u> file on the GitHub repository for information about ways the ACDP can be configured for public internet access, private network only access, and access with Transit Gateway.

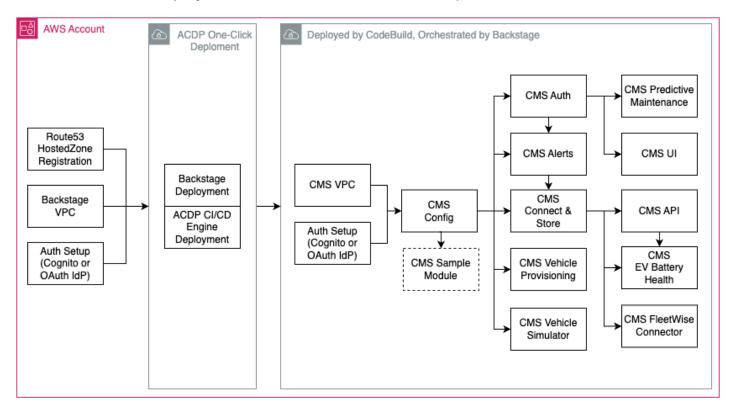
Deploying CMS on AWS modules with Backstage

The AWS CodePipeline deployment of Backstage initiates from the ACDP deployment described above allowing for Backstage to deploy additional CMS on AWS modules. Deploying CMS on AWS modules using Backstage is represented by the following diagram:



13

All CMS on AWS Modules are deployable through <u>Backstage templates</u>. These deployments are powered by AWS CodeBuild. Some modules have deploy-time dependency on other modules and therefore should be deployed in the order of those relationships (shown below).



CMS on AWS modules and services

The high-level architectural descriptions for the CMS on AWS modules and services are as follows:

Amazon Virtual Private Cloud (Amazon VPC)

Amazon Virtual Private Cloud (Amazon VPC) is an AWS service that allows you to launch AWS resources inside a logically isolated virtual network. CMS on AWS provides a VPC module that deploys an opinionated network configuration. For more details, see <u>Virtual Private Cloud</u> and <u>Amazon VPC</u>.

Auth Setup

The Auth Setup module provides the means to configure a third-party OAuth 2.0 compliant IdP of your choice for use with CMS on AWS. An optional deployment of Amazon Cognito infrastructure is also provided, with pre-populated configurations for integrating with the rest of CMS on AWS. The

Auth Setup module either deploys configurable secrets with a defined JSON structure, or it can use existing Secrets Manager secrets. Either will fully configure the parameters required for a CMS on AWS deployment's authentication. For more information, see Auth Setup.

Config

The CMS Config module uses the <u>AWS Systems Manager Parameter Store</u> to register a unique ID which serves as a namespace to deploy other CMS on AWS modules. The CMS Config module takes the VPC name and Identity Provider ID as additional inputs which are shared with the other CMS on AWS modules deployed with the same unique ID as the CMS Config module. The module uses an <u>AWS Lambda</u> function to send anonymized metrics about AWS S3 and AWS Timestream resource usage. The module also implements an AWS Lambda function for AWS SSM Parameter resource lookup based on the unique ID.

Automotive Cloud Developer Portal (ACDP) and Backstage

The Automotive Cloud Developer Portal (ACDP) is the centralized platform for deploying subsequent CMS on AWS modules. The ACDP uses the Backstage module as its presentation layer to provide a configurable developer platform for managing and monitoring the deployment of CMS on AWS modules and customer provided modules. ACDP also supports cross-account and cross-region deployments of modules. For more details, see Automotive Cloud Developer Portal and Backstage module.

Auth

The CMS Auth module allows for the authentication and authorization of users and services throughout the solution. The module provides two AWS Lambda functions which can integrate with any third-party identity provider (IdP) that is OAuth2.0 OIDC compliant. This is done by communicating with the <u>Auth Setup</u> module's IdP configurations, exposed as Secrets Manager secrets. Of these two Lambda functions, one facilitates exchanging an authorization code for an access token with the <u>authorization code flow</u>, and the other validates and authorizes access tokens. For more details, see <u>Auth module</u>.

AWS IoT Core and MQTT

AWS IoT Core MQTT topics are the primary method for communicating events between the CMS on AWS modules. Messages published to MQTT from CMS on AWS modules can be consumed by, and invoke rules configured by, other modules. AWS IoT Core is also used as the primary management

Config 14

and storage system for provisioned vehicles. For more details on the usage of AWS IoT Core, see Vehicle Provisioning module.

Vehicle Provisioning

The CMS Vehicle Provisioning module provides means to onboard and register vehicles with AWS IoT Core. Deploying the module checks for the existence of, and if not found creates, a claim certificate and private key pair for use with fleet provisioning by claim. This claim certificate is linked to a well-defined <u>provisioning template</u>, which controls how vehicles are provisioned and informs the <u>AWS IoT policy</u> that is given to newly provisioned vehicles.

Using the claim certificate, a vehicle can retrieve a unique certificate to allow for further communication with AWS IoT Core. Registering invokes <u>AWS IoT rules</u> linked to Lambda functions. These functions check for vehicle authorization and create and manage vehicle records in <u>Amazon DynamoDB</u>. At the end of the process, the solution registers an AWS IoT Core thing for the vehicle that is linked to credentials safely stored in <u>AWS Secrets Manager</u>. For more details, see <u>Vehicle Provisioning module</u>.

Connect and Store

A centralized <u>Amazon S3</u> bucket deployed within the Connect and Store module serves as the reservoir for all CMS on AWS data. Centralized data storage allows for querying of vehicle telemetry data and enabling alerts based on data insertion and thresholds. For more details, see Connect and Store module.

FleetWise Connector

The CMS FleetWise Connector module allows you to consume data that is captured by <u>AWS IoT FleetWise</u> campaigns. This is done by querying Amazon Timestream to migrate data into the CMS on AWS Connect & Store module's telemetry bucket. The data is then indexed using <u>AWS Glue</u>, and made accessible through <u>Amazon Athena</u>.

API

CMS on AWS users can interact with vehicle telemetry data stored in the CMS on AWS data lake through the CMS API module. API endpoints are provided through <u>AWS AppSync</u>, which expects GraphQL requests. AWS AppSync endpoints use Lambda functions to build and run <u>Amazon</u> Athena queries on vehicle data stored in Amazon S3.

Vehicle Provisioning 15

Alerts

The CMS Alerts module allows you to receive notifications invoked by data stored in the CMS on AWS data lake. CMS modules can publish to <u>Amazon SNS</u> topics defined by the CMS Alerts module by utilizing an API provided through AWS AppSync. You can subscribe to these same topics to receive email notifications. For more details, see <u>Alerts module</u>.

EV Battery Health

For monitoring stored data, CMS on AWS users can use the CMS EV Battery Health module. This module provides a dashboard through <u>Amazon Managed Grafana</u>, which is authenticated by <u>AWS IAM Identity Center</u>. From the dashboard, users can visualize data and setup alerts based on configurable data thresholds. For more details, see EV Battery Health module.

Vehicle Simulator

The CMS Vehicle Simulator module provides a user interface (UI) and backend engine for creating, operating, and monitoring simulations of vehicle data emissions. Simulations are configurable by interval, number of vehicles, and overall durations. They also support either a custom payload schema, or the provided default VSS schema.

This solution runs simulations by using <u>AWS Step Functions</u>, backed by a series of AWS Lambda functions. The simulator handles scaling, AWS IoT Core provisioning and registration, telemetry data generation, generation intervals, and total emission quantity. For more details, see <u>Vehicle Simulator module</u>.

Predictive Maintenance

The CMS Predictive Maintenance module provides an MLOps infrastructure implemented using Amazon SageMaker AI Pipelines. The pipeline implements pre-processing, training, evaluating, model accuracy checking, and model deployment stages. You can use the deployed model to perform inference by using Amazon SageMaker AI Serverless Inference for real-time inference and batch transform for batch inference. The module exposes two APIs by using Amazon API Gateway for performing real-time inference and batch inference.

The module additionally implements a generative AI chatbot using by using an <u>Amazon Bedrock Agent</u> to obtain fleet insights using natural language queries. The chatbot uses <u>Amazon Bedrock Knowledge Bases</u> to perform Retrieval-Augmented Generation (RAG) and uses action groups to

Alerts 16

query the backend by using Lambda functions to obtain real-time vehicle details. For more details, see Predictive Maintenance module.

Fleet Management UI (Preview)

The CMS Fleet Management UI module provides a user interface (UI) and backend middleware that integrates with <u>AWS IoT FleetWise</u> to enable creation of fleets and vehicles. It contains multiple demo dashboards to showcase the art-of-the-possible for how data collected by <u>AWS IoT FleetWise</u> could be used to monitor vehicle and fleet metrics.

The module is currently in a preview state to showcase capabilities that can be enabled via integration with <u>AWS IoT FleetWise</u>. For more details, see <u>Fleet Management UI module</u>.

AWS Well-Architected design considerations

This solution uses the best practices from the <u>AWS Well-Architected Framework</u>, which helps customers design and operate reliable, secure, efficient, and cost-effective workloads in the cloud.

This section describes how the design principles and best practices of the Well-Architected Framework were applied when building this solution.

Operational excellence

This section describes how we architected this solution using the principles and best practices of the operational excellence pillar.

The built-in CI/CD pipeline enables a standardized deployment strategy for the ACDP and Backstage module, as well as supporting the further managed deployment of CMS on AWS modules with AWS CodeBuild. CMS on AWS sends logging and metrics to CloudWatch throughout the entire solution. A default log retention of three months is used in most places; this can be customized by altering the CDK (look for aws_logs.RetentionDays) and rebuilding the solution. The infrastructure is managed and operated by AWS CDK, with deployment assets stored in Amazon S3 for use with Backstage.

Security

This section describes how we architected this solution using the principles and best practices of the security pillar.

To ensure network security, CMS on AWS network traffic only flows through the internet when necessary. The traffic flows between AWS services through the VPC network and VPC endpoints when possible. Simultaneously, all internet accessible endpoints are protected by authentication (OAuth2.0) through the customer chosen identity provider, expecting the configured issuer, audience, clients, and scopes with access tokens. All databases are closed off from anything external to the AWS account where they are deployed and have rotating credentials where applicable. CMS on AWS data is encrypted at rest and in transit with rotating encryption keys. Permissions are locked down to zero-trust principles or least-privilege; the most restrictive choice is made where possible. Parameters that contain sensitive information are stored in AWS Secrets Manager and rotated, while remaining parameters are kept in Parameter Store, a capability of AWS Systems Manager.

This solution's default configuration doesn't deploy a web application firewall (WAF) in front of API endpoints. To enhance your API security with a WAF, you must set it up manually. AWS provides an in-depth guide on how you can control access to API Gateway with AWS WAF. For instructions on how to implement AWS WAF in front of your API and increase distributed denial of service (DDoS) protection for your web applications, see Setting up AWS WAF and its components.

Reliability

This section describes how we architected this solution using the principles and best practices of the <u>reliability pillar</u>.

CMS on AWS uses primarily serverless AWS services (a notable exception being Backstage), which provides resiliency, uptime, and automatic scaling. All appropriate Amazon S3 buckets have versioning enabled and are backup protected. All DynamoDB tables have point-in-time recovery, and customer data is not deleted when you uninstall the solution.

Performance efficiency

This section describes how we architected this solution using the principles and best practices of the <u>performance efficiency pillar</u>.

All compute and performance efficiency relates to usage and not a base cost. Complex tasks are delegated to appropriate AWS services that provide built-in, efficient functionality to minimize needed compute resources and prevent bottlenecks. You can deploy in <u>supported AWS Regions</u> to keep your data closer to where it's being used and processed, minimizing delays.

Reliability 18

Cost optimization

This section describes how we architected this solution using the principles and best practices of the cost optimization pillar.

<u>AWS Billing and Cost Management</u> provide cost observation and analysis. CMS on AWS follows a consumption model, so costs are driven by usage.

You can also view the total cost from the **Metrics** tab in Backstage after deploying a module. You can find further cost breakdowns in the myApplications dashboard.

Sustainability

This section describes how we architected this solution using the principles and best practices of the sustainability pillar.

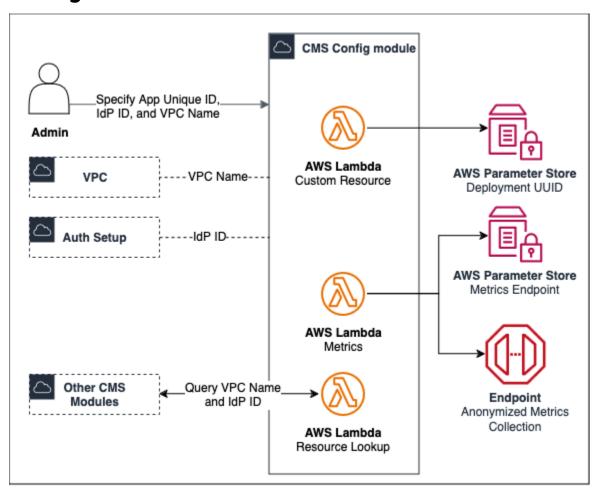
This solution uses primarily managed and serverless services to minimize the environmental impact of the backend services.

Cost optimization 19

Architecture details

This section describes the components and AWS services that make up this solution and the architecture details on how these components work together.

Config



The Config module is a prerequisite deployment required for the deployment of any CMS on AWS module (all modules except ACDP, VPC, and Auth Setup). The Config module introduces the concept of an App Unique ID, which serves as a unique namespace for all CMS on AWS modules. This module enables sharing the same VPC and IdP among multiple CMS on AWS modules without requiring input of the VPC and IdP configuration in each module's deployment. The App Unique ID serves the following purposes:

• A unique naming prefix for all AWS CloudFormation resources that avoids conflicts when deploying multiple instances of the same CMS on AWS module in the same region and account

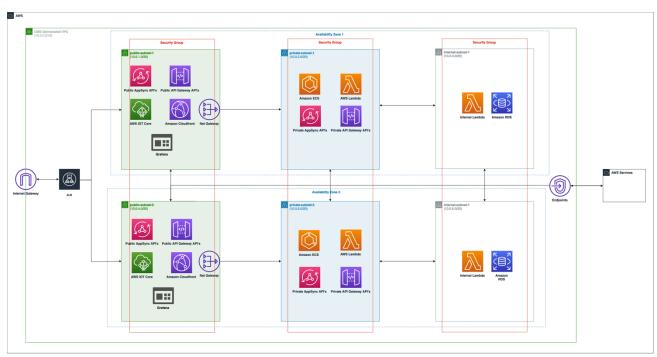
Config 20

- A unique naming prefix for easy lookup of AWS SSM Parameters
- Registering deployment of CMS on AWS modules associated with the App Unique ID, preventing multiple deployments of the same module within the same CMS on AWS deployment

The CMS Config module also provides AWS Lambda functions that perform the following functions:

- Create a deployment UUID which is used to tag all CloudFormation resources created by CMS on AWS modules
- Send anonymized metrics related to Amazon S3 and AWS Timestream
- · Lookup VPC name and Identity Provider ID using the App Unique ID

Virtual Private Cloud (VPC)



CMS on AWS requires a VPC to deploy other modules including ACDP. Users can bring their own VPC or deploy the CMS on AWS VPC module. The VPC module provides a preconfigured VPC with the architecture shown in the diagram. It has the following configuration:

- Public VPC:
 - Two Availability Zones (AZ)
 - (Optional) VPC endpoints for most of the AWS services that support AWS PrivateLink

Virtual Private Cloud (VPC) 21

- One NAT gateway in each AZ
- Public, private, and isolated subnets in each AZ
- One internet gateway
- Private VPC:
 - Two Availability Zones (AZ)
 - (Optional) VPC endpoints for most of the AWS services that support AWS PrivateLink
 - Attachable, private, and isolated subnets in each AZ
 - (Optional) Transit Gateway attachment

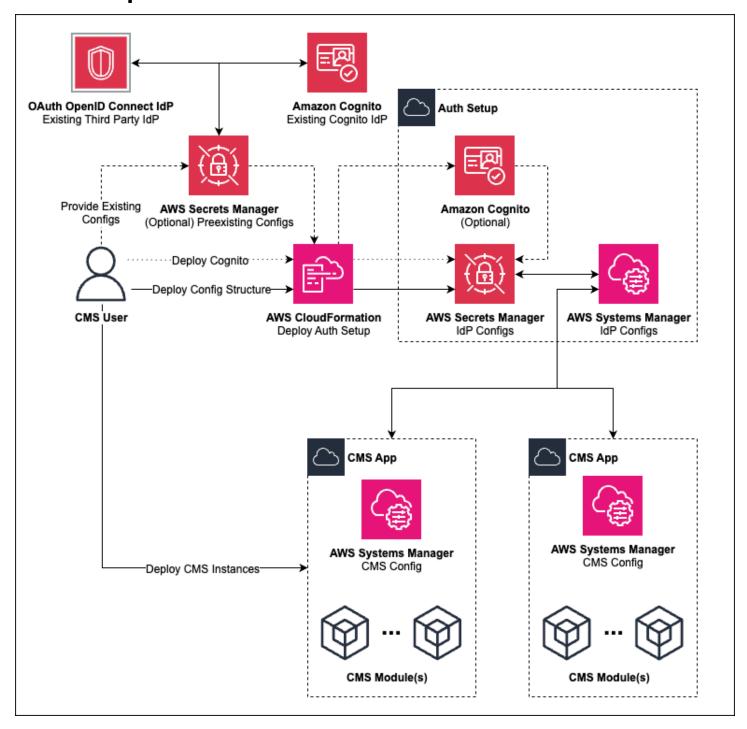
Two AZs and two NAT Gateways ensures that all the services are still functioning in case there is a failure in one of the two AZs. The three subnets in each AZ are for the modules to allocate resources as needed.

- Public/attachable subnet for resources that require inbound and outbound connection to the internet through internet gateway or transit gateway
- Private subnet for resources that require an outbound only connection to the internet through NAT gateway
- Isolated subnet for resources that do not require any inbound or outbound connection to the internet

Resources should be assigned a security group at the module level depending on their needs.

Virtual Private Cloud (VPC) 22

Auth Setup



The Auth Setup module creates Secrets Manager secrets, which provide necessary configurations for other CMS on AWS modules. This includes the CMS Auth module's ability to communicate with the IdP and integrating the Auth Setup module's IdP for use with Backstage.

Auth Setup 23

The configurations include identity provider details including, but not limited to, relevant endpoints, expected domains, audiences, scopes, and client identifiers.

This module also provides an optional deployment of Amazon Cognito infrastructure that is preconfigured to serve as the IdP for the CMS on AWS deployment, including authentication for Backstage.

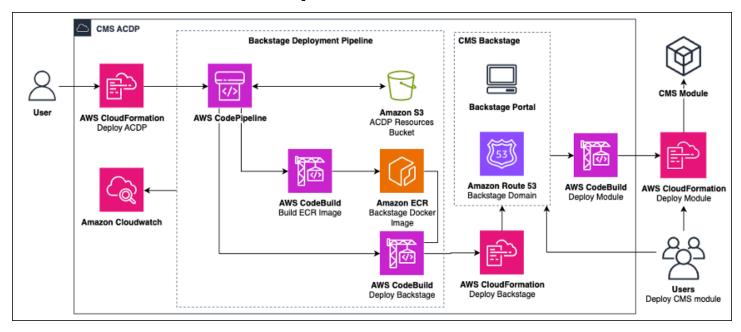
Each configuration secret is a known JSON structure with consistent key/value pairs expected by CMS on AWS modules. The key structure is consistent regardless of deployment path, but the values are populated dynamically. The three deployment paths are as follows:

- Deploy Amazon Cognito infrastructure If choosing to deploy Amazon Cognito infrastructure, the configuration JSONs will be pre-configured with values from the newly deployed Amazon Cognito resources.
- **Provide existing configs** During deployment, you are given the option to provide ARNs for zero to all of the configuration secrets. If providing an existing secret ARN, a new secret will not be created and no validation of the secret value's structure will be performed. The existing secret will be configured to be used by the CMS on AWS deployment.
- Deploy config structure During deployment, you are given the option to provide ARNs for zero to all of the configuration secrets. If not providing an existing secret ARN, a new secret will be created with the expected JSON structure. The values of this JSON structure will be empty however and will need to be manually set for the CMS on AWS deployment's authentication to function properly.

Since the choice to provide an existing secret ARN is individual to each config, the latter two deployment paths can be combined in a single deployment for differing secrets.

Auth Setup 24

Automotive Cloud Developer Portal



The initial deployment of CMS on AWS includes the ACDP and Backstage, which assist in managing the deployment of CMS on AWS modules. The ACDP provides a presentation layer through the Backstage module, which you can use to select, configure, and deploy individual modules. Backstage also allows for monitoring and teardown of these modules, as well as viewing of module documentation. Backstage integrates with AWS CodeBuild for deployment execution. The ACDP is a powerful platform that enables flexibility and quick insights into deployed infrastructure.

Deploying the ACDP creates and configures an AWS CodePipeline for deploying the Backstage module. The pipeline is configured with the necessary environment to carry out the Backstage deployment. Values for this configuration are taken from the local environment of the ACDP deploy, which is configured with Make. An Amazon ECR repository is created and permissions are shared between it and the Backstage pipeline.

The pipeline then handles the deployment of the Backstage module in three stages. These stages are abstracted from the previous diagram:

- 1. **Source** Sources the Backstage source code from an Amazon S3 bucket.
- 2. **Build** Builds the Backstage docker image and pushes it to the private Amazon ECR repository.
- 3. **Deploy** Deploys the Backstage module stack.

Once deployed, you can then use the Backstage module to deploy CMS on AWS modules.

Supported networking scenarios for ACDP

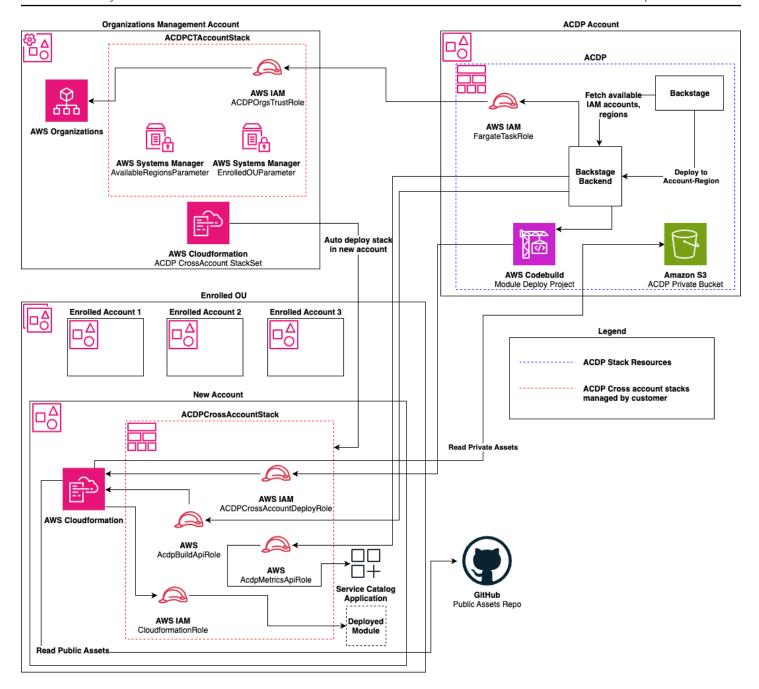
ACDP can be configured with the following networking topologies:

- 1. Publicly accessible through the internet with Route53 and an auto generated ACM certificate.
- 2. Publicly accessible through the Internet with an external DNS and a custom certificate uploaded to ACM.
- 3. Private endpoints with public DNS hosted by Route53 and an auto generated ACM certificate.
- 4. Private Route53 Hosted Zone with a custom certificate uploaded to ACM.
- 5. Private VPC linked to a public VPC through a transit gateway.

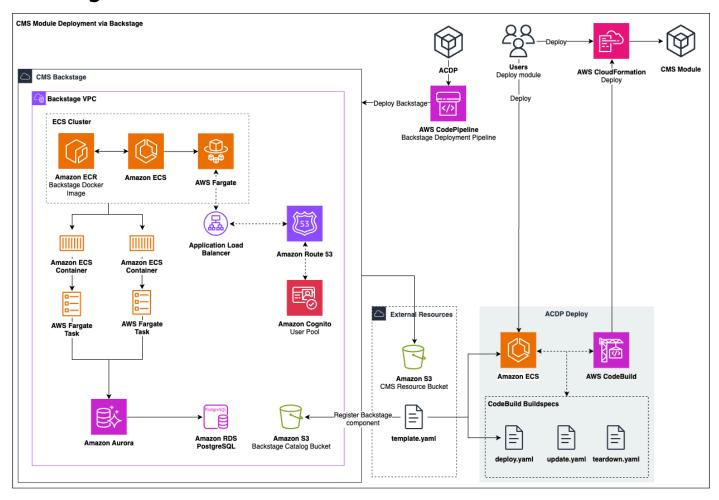
Multi Account/Multi Region Deployment

ACDP can be configured to enable cross-account deployment, thus enabling a Multi Account/ Multi Region Deployment feature. In order to enable this feature, you need to provide values for certain CloudFormation Parameters when you deploy ACDP CloudFormation template. Detailed instructions are given in Deployment Instructions section.

Upon succesfull deployment of ACDP, you need to deploy the <u>Guidance stack</u>. This stack deploys SSM Parameters necessary to discover enrolled accounts and regions, as well as a StackSet configured to deploy AWS IAM roles necessary for ACDP to do cross-account deployments. This StackSet is automatically triggered when a new account is added to the enrolled Organizational Unit.



Backstage module



CMS on AWS uses Backstage as the preferred approach for deploying its modules. The Backstage module is deployed and configured to deploy within the provided <u>VPC</u> and requires an Amazon Route53 hosted zone. This hosted zone is specified in the deployment parameters as the host for the Backstage portal (see Working with hosted zones).

The Backstage module also provides both authentication and authorization by default. Authentication is accomplished by communicating with the identity provider configured via the Auth Setup module. Authenticated users can then log in to the Backstage portal. Authorization is implemented into Backstage via the Backstage portal. Authorization is configured to deny by default. All permissions can be managed via the RBAC front-end, where they can be grouped into roles and assigned to users. On initial deployment, only a single super user is created, which will be the starting point for granting further permissions to other users.

The Backstage module is deployed as part of the ACDP deployment. The ACDP deployment contains two main parts: the AWS Infrastructure required by the Backstage deployment pipeline,

Backstage module 28

and a pipeline to build and deploy a copy of Backstage in the specified account. The infrastructure deployed by the pipeline includes an Amazon ECS Fargate job to run Backstage's Docker image and Aurora PostgreSQL-Compatible Edition database, an Amazon S3 bucket that acts as the Backstage catalog, and the Amazon Route 53 domain setup. This infrastructure makes up the Backstage module. Authentication is configured for and handled by the Backstage API and therefore requires an existing identity provider and its corresponding configurations deployed by the Auth Setup module.

Backstage uses an Amazon VPC (Virtual Private Cloud) to create a private network to help protect some of its resources. Specifically, Backstage uses the following resources within its VPC:

- **ELB** The Application Load Balancer connects with the Amazon Route 53 domain, as well as the Backstage Fargate service, to help orchestrate and balance tasks.
- Amazon ECS Creates a cluster, combined with AWS Fargate, to group task definitions and provide the container image through Amazon ECR.
- Amazon ECR Stores the Backstage image to be supplied to the docker container associated with the AWS Fargate task definitions.
- AWS Fargate Combined with Amazon ECS, allows for defining task definitions with associated containers, and running those containers without needing to manage Amazon EC2 instances.
- Amazon Aurora PostgreSQL The relational database used by Backstage.

As mentioned previously, authentication to the Backstage portal is handled by using the sign-in functionality provided by the Backstage API. Enabling this functionality requires configuring the identity provider of choice, which is done through the <u>Auth Setup</u> module. Backstage implements a popup flow by default, but an experimental redirect flow is also offered and is required for integrating Backstage with Amazon Cognito. It is also possible to customize the scopes which Backstage will request during its authentication attempt. This is important for properly configuring with some identity providers, and also allows for bespoke implementation of authorization based on authenticated user access.

The Backstage module is integrated with AWS CodeBuild to enable its deployment functionality. The module deploys with pre-registered Backstage templates for each of the CMS on AWS modules. Deploying a CMS on AWS module through the Backstage UI requires providing required parameters, and in some cases, deploying the modules in the right order as shown in the deployment diagram. When a Backstage component is deployed, it uses the template. yaml file to instruct the process and performs the following steps:

Backstage module 29

- 1. Copy required deployment assets and docs to the Amazon S3 catalog bucket.
- 2. Write the module's catalog info to the Amazon S3 catalog bucket.
- 3. Register the module within the Backstage module's catalog.
- 4. Configure the ACDP deployment backend to be able to deploy the module.

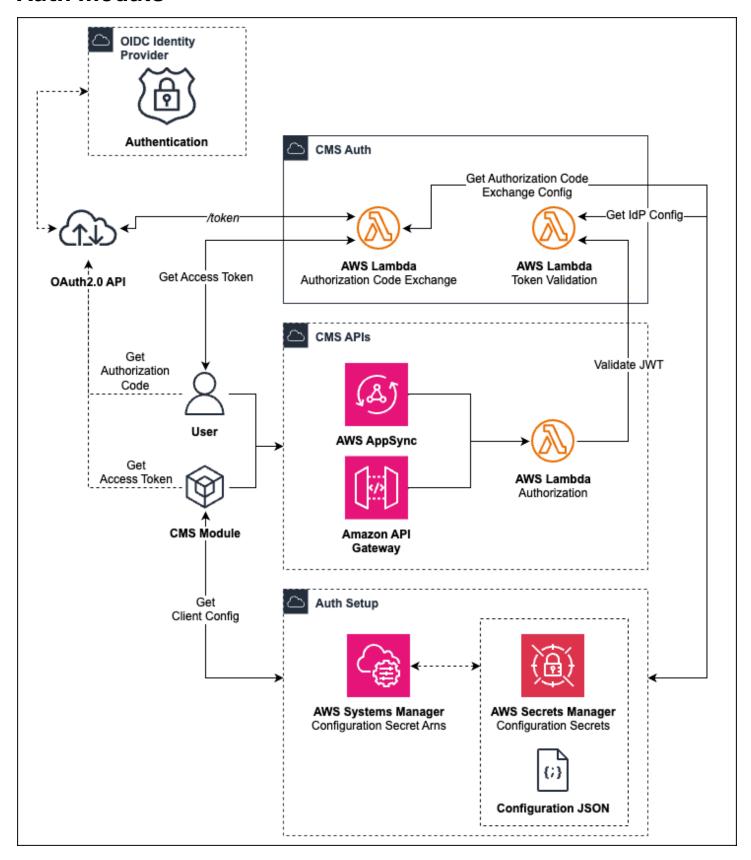
After creating the catalog-info.yaml, deployment progress can be monitored on the **CI/CD** tab of the Backstage Catalog Item, which provides visibility into the CodeBuild results for your deployment.

Similarly, you can see the cost of your deployment in the **Metrics** tab of the Backstage Catalog Item. You can also use the myApplications dashboard for further monitoring of your deployment.

Lastly, the Backstage module can be used to view documentation related to each module directly from within the portal. These docs are included in the assets that are built and uploaded to Amazon S3 for use by Backstage.

Backstage module 30

Auth module



Auth module 31

The Auth module provides the means for CMS on AWS users and services to authenticate and authorize themselves for use with CMS on AWS APIs and portals. Users signed up with the identity provider can request authorization codes or access tokens from the authorization server with authorization code flow, or implicit flow. Services can request access tokens from the authorization server with client-credentials flow. In either case, the access token will then be validated by the identity provider resource server to grant access to CMS on AWS APIs. Access tokens can also be used to grant access to be poke implementations of CMS on AWS front-end portals.

These authentication flows and functionality are supported by two Lambda functions deployed by the Auth module: the token validation and token exchange Lambda functions.

The token validation Lambda function implements a JWT validation flow that is defined by the standard OAuth2.0 protocol This validation can be done for both user and service access tokens and is performed as follows:

- 1. Validate the integrity of the JWT signature.
- 2. Check the token expiration.
- 3. Authorize the token's claims and scope against the identity provider configuration provided by the Auth Setup module.



Note

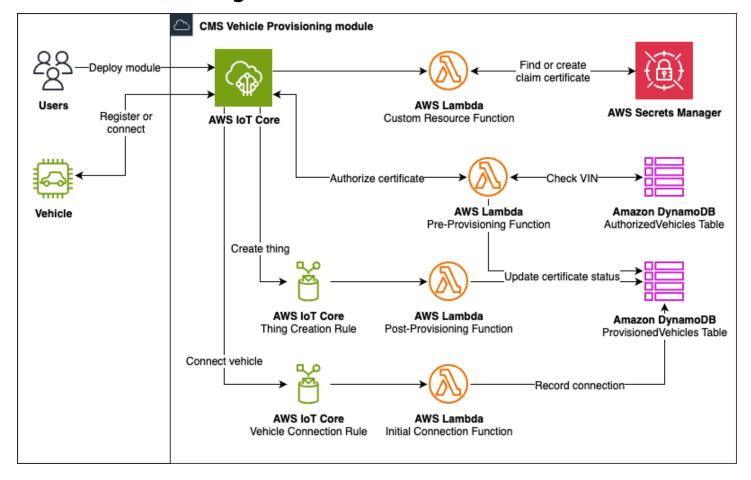
The details of this process are abstracted from the diagram above.

The token exchange Lambda function implements an authorization code exchange, defined by the authorization code flow, to retrieve an access token from the /token endpoint. The access tokens can then be used to authenticate and authorize users against CMS on AWS APIs. The token exchange Lambda function communicates with the authorization server, specified in the identity provider config secret, to verify the integrity of the authorization code against the associated identity provider.

The token exchange Lambda function allows for usage of an optional Proof Key for Code Exchange (PKCE) code verifier to protect against injection attacks, which could intercept the access token.

Auth module 32

Vehicle Provisioning module



The Vehicle Provisioning module defines, creates, and manages the certificates, policies, and vehicle profiles for registered vehicles. This module allows vehicles to register with AWS IoT Core using a secure communication system encrypted with the <u>TLS v1.2 protocol</u>. On deployment, the module generates a unique claim certificate to allow vehicle registration. Registered vehicles receive credentials to allow connections to AWS IoT Core.

This module uses the fleet provisioning by claim workflow which is supported by AWS IoT Core. When the CloudFormation stack is created or updated, a custom resource Lambda function initiates, configuring AWS IoT Core to enable AWS IoT Core thing events for detecting vehicle registrations. The custom resource function retrieves an existing claim certificate and private key from AWS Secrets Manager. If not found, the claim certificate and private key are created, activated, and stored in AWS Secrets Manager. The claim certificate has an attached provisioning template that configures AWS IoT Core thing, certificate, and policy creation.

Vehicle Provisioning module 33

After setup, a user can connect to AWS IoT Core using the claim certificate, private key, and Amazon Root CA, which is used to sign the claim certificate. After connecting, the user can retrieve a new unique certificate and private key for registering the vehicle.

To register the vehicle, the solution calls the AWS IoT Core RegisterThing endpoint with credentials. Starting the registration process invokes the <u>pre-provision Lambda function</u>, which completes the following:

- 1. Finds existing certificates for this vehicle in the ProvisionedVehicles DynamoDB table. If certificates are found and not already INACTIVE, the function deactivates them and updates the ProvisionedVehicles records to reflect the change.
- 2. Creates a record in the ProvisionedVehicles DynamoDB table for the new combination of Vehicle Identification Number (VIN) and certificate in the PENDING_ACTIVATION status.
- 3. Searches for the vehicle in the AuthorizedVehicles DynamoDB table. If the vehicle is not found, the solution prevents registration by deleting the certificate.
- 4. Returns a registration-allowed Boolean.

If provisioning was allowed for this vehicle, AWS IoT Core creates or updates the thing for this vehicle. The certificate for this vehicle is activated, and a policy is created and attached to the certificate. This policy is defined in the provisioning template.

Continuing the registration process creates an AWS IoT Core thing that invokes a post-provision Lambda function, which completes the following:

- 1. Updates the new certificate's status to ACTIVE in the ProvisionedVehicles DynamoDB table.
- 2. Deletes this vehicle's INACTIVE certificates from AWS IoT Core.
- 3. Deletes certificates' record in the ProvisionedVehicles DynamoDB table.

After a registered vehicle connects to AWS IoT Core, the vehicle can then publish to the vehicleactive AWS IoT MQTT topic to signal a successful connection to the solution. Messages to this topic invoke the initial-connection Lambda function, which flips a Boolean for the vehicle record to indicate that the vehicle has successfully connected with their certificate at least one time.

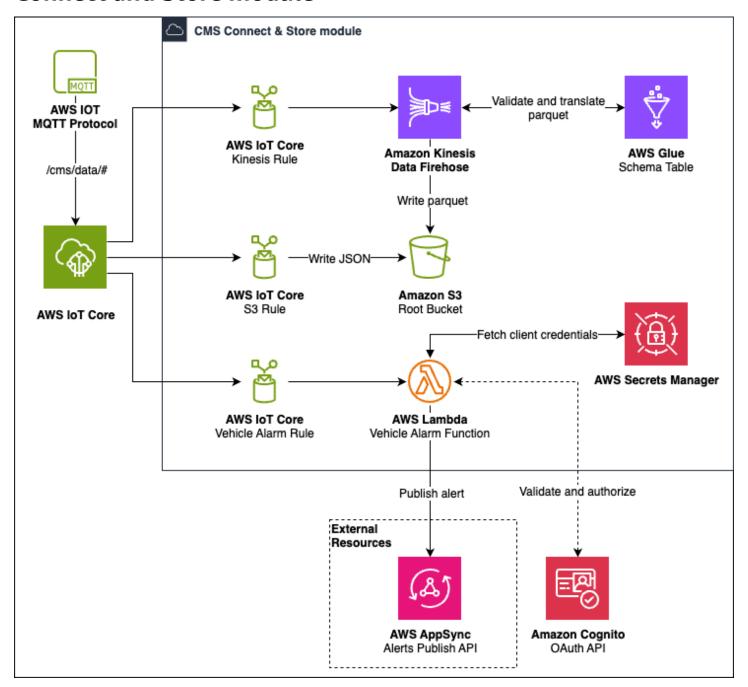
Vehicle Provisioning module 34



Note

Publishing to the vehicleactive topic is not an automatic part of registration, and is not implemented as a part of the solution.

Connect and Store module



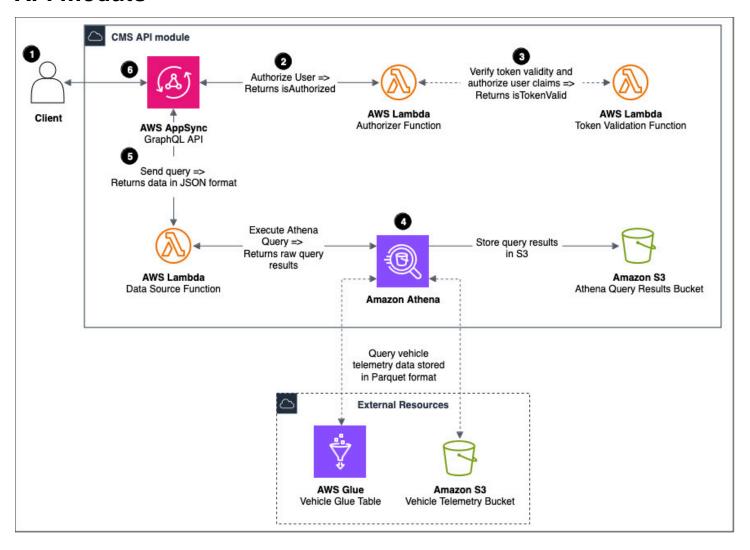
Connect and Store module 35 The Connect and Store module is the primary data lake for the solution. This module provides the collection of telemetry data payloads from well-defined AWS IoT MQTT topics. Data is stored in an Amazon S3 bucket in both JSON and Apache Parquet format. Centralizing all CMS on AWS telemetry data into a single data lake enables the CMS on AWS modules to retrieve data without needing to interface with AWS IoT Core directly.

A single Amazon S3 bucket acts as the data lake within the module. The Amazon S3 objects are prefixed with the timestamp, data format, and vehicle identifier.

This module uses three AWS IoT rules. Two of those rules subscribe to a broad AWS IoT MQTT topic invoked on vehicle data ingestion. The first rule invokes an action property that writes the JSON payload to the appropriate Amazon S3 bucket. The second rule invokes an action property that passes the payload to Amazon Data Firehose. A delivery stream validates and transforms the payload using AWS Glue. The stream then writes the Parquet-formatted payload to the appropriate Amazon S3 bucket. Using Firehose allows the Connect and Store module to handle a high throughput of telemetry data payloads. The third rule subscribes to a broad MQTT topic, which can be invoked by vehicle emission. This rule delivers the received notification to the intended CMS Alerts API, which forwards the notification to the appropriate destination.

Connect and Store module 36

API module



The API module provides the ability for users and services to query data stored by the solution. This is done by leveraging an AWS AppSync GraphQL API that integrates with the default CMS on AWS data lake deployed from the Connect and Store module. You can integrate it with a data lake or resources external to the solution or pre-built modules by changing the module's configuration.

The AWS AppSync API uses Amazon Athena with a Lambda resolver data source, allowing the data to be queried the moment it is stored in the solution. This provides a near real-time representation of the state of vehicles.

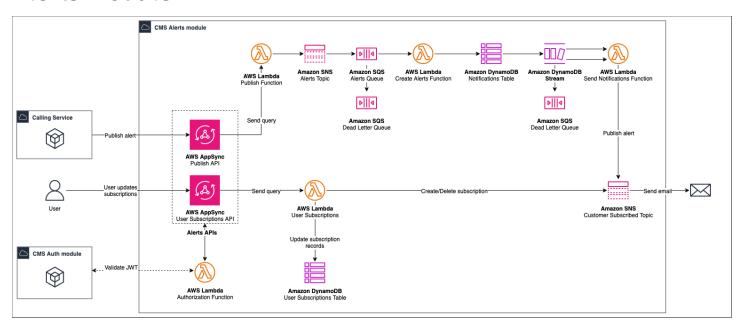
By default, the CMS API module integrates with the Connect and Store module's Amazon S3 bucket as its data source and the CMS Auth module's IdP for authorization on API requests.

The following steps explain how a request to the API is handled:

API module 37

- 1. The client makes a valid GraphQL request to the API endpoint with a bearer token provided in the authorization header. The bearer token should be an access token obtained from the token endpoint of the configured IdP.
- 2. The token must be validated and authorized against the chosen IdP's user pool (this functionality is available through the CMS Auth module). Further authorization logic can determine whether the user has permission for the operations and fields selected.
- 3. Once authorized, the context of the GraphQL query is sent to a Lambda resolver. The resolver builds and invokes an Amazon Athena query using the selected fields and provided arguments.
- 4. Amazon Athena uses the configured AWS Glue table to query the vehicle data from the data lake, and store the results in a separate Amazon S3 bucket.
- 5. The Lambda resolver parses the results into JSON format and returns them to AWS AppSync.
- 6. The AWS AppSync API receives the results and returns them to the client.

Alerts module



The Alerts module enables CMS on AWS and customer implemented modules to send alerts to subscribed users, and allows users to manage their alert subscriptions. This is done by leveraging two AWS AppSynchttps://graphql.org/learn/[GraphQL] API operations; one for user subscription management and another for publishing messages to user subscribed Amazon SNS topics. The CMS Alerts module requires the CMS Auth module as a prerequisite to authenticate API requests.

Alerts module 38

The user subscription AWS AppSync API uses <u>Amazon DynamoDB</u> with a Lambda resolver data source. This Lambda function stores user subscription information in a DynamoDB table as well as subscribes and unsubscribes users from an Amazon SNS topic.

The following steps explain how a request to the user subscription API is handled (left to right in diagram):

- 1. The client makes a valid GraphQL request to the API endpoint with a bearer token provided in the authorization header. The bearer token should be an access token obtained from the token endpoint of the configured IdP.
- 2. The token must be validated and authorized against the chosen IdP's user pool (this functionality is available through the CMS Auth module). Further authorization logic can determine whether the user has permission for the operations and fields selected.
- 3. Once authorized, the context of the GraphQL query is sent to a Lambda resolver. The resolver updates the user's subscription preferences on Amazon SNS and makes a corresponding update to the DynamoDB table.
- 4. The Lambda resolver parses the results into JSON format and returns them to AWS AppSync.
- 5. The AWS AppSync API receives the results and returns them to the client.

The publish API endpoint uses a Lambda function as its data source. When a CMS on AWS module sends a request, the Lambda function relays the message through the system.

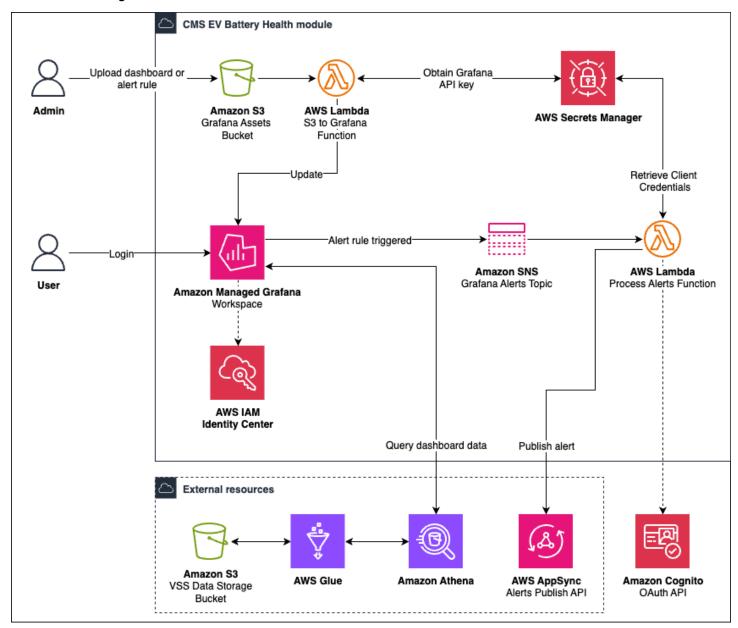
The following steps explain how a request to the publish API is handled (left to right in diagram):

- 1. The client makes a valid GraphQL request to the API endpoint with a bearer token provided in the authorization header. The bearer token should be an access token obtained from the token endpoint of the configured IdP.
- 2. The token must be validated and authorized against the chosen IdP's user pool (this functionality is available through the CMS Auth module). Further authorization logic can determine whether the user has permission for the operations and fields selected.
- 3. Once authorized, the context of the GraphQL query is sent to a Lambda resolver. The Lambda resolver publishes this message to the central Amazon SNS topic, which notifies the central Amazon Simple Queue Service (Amazon SQS) queue.
- 4. This queue triggers the create-alerts Lambda function, which stores this message in a notifications DynamoDB table.

Alerts module 39

- 5. This DynamoDB table has a <u>DynamoDB stream</u> enabled, which notifies a send-notifications Lambda function about the changes in the DynamoDB table.
- 6. The send-notifications Lambda function publishes all the notifications in the stream to their corresponding Amazon SNS topics.
- 7. The users subscribed to these Amazon SNS topics receive an email notification.

EV Battery Health module



The EV Battery Health module implements an Amazon Managed Grafana workspace to facilitate visualizing and alerting based on vehicle telemetry data. The workspace is created with AWS IAM

EV Battery Health module 40

Identity Center authentication. The workspace uses Amazon Athena as its query engine for panels and alerts, with queries written in <u>ANSI SQL</u>. The Amazon Managed Grafana console is used to add users to the workspace and assign users to Admin, Editor, or Viewer roles.

After deployment, the solution creates an Admin API key for the workspace, which allows creating and updating the dashboard panels and alert rules. Amazon Managed Grafana API keys expire in 30 days, so the API key is stored in Secrets Manager and rotated every 29 days. The workspace is provisioned with unified alerting to leverage the alerting capabilities of Amazon Managed Grafana.

The EV battery health dashboard contains gauge panels for Remaining Useful Life (%), Remaining Charge (%), Average Temperature (C°), and Current Voltage (V). The solution creates stat panels to visualize the remaining driving range from battery and hybrid energy sources. The solution also creates a time series graph panel, which shows the historical data for the Remaining Useful Life (%) of the battery. The dashboard can be parameterized by VIN, allowing data to be visualized for each vehicle by selecting its corresponding VIN from the dropdown at the top of the dashboard. The dashboard data model is stored in a versioned and encrypted Amazon S3 bucket, which uses an event-driven approach to invoke a Lambda function if the data model object is modified. This facilitates modifying the dashboard data model post deployment without having to update the deployment itself.

The solution creates alert rules to send an alert message whenever the remaining charge or remaining useful life of the battery drops below a configurable threshold. A time series query is run, and the results of the query are reduced to a single value by a reducer function. The reduced value is then compared with a threshold and evaluated for whether the queried data violates the alert rule.

You can configure alert rules as to how often they evaluate queries, and how long they should wait before firing an alert. If the alert rule is breached, the solution sends an alert message to an Amazon SNS topic. The Amazon SNS topic has a Lambda function subscriber that processes the alert messages. The solution sends messages to the alerts endpoint exposed by the CMS Alerts module. The API call to the CMS Alerts module is authenticated by an access token obtained through the client credentials flow, using the service app client provided by the CMS Auth module.

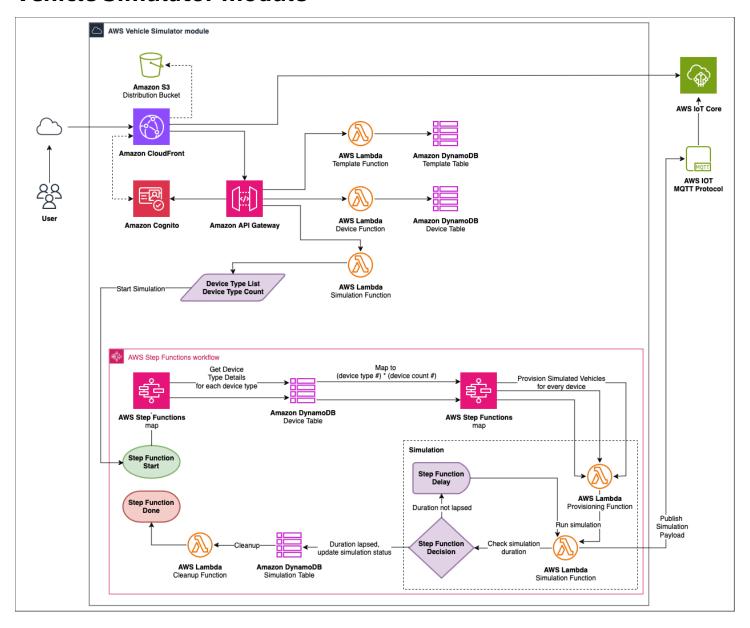
The deployment of the CMS EV Battery Health module uses the AWS CDK library to create the CloudFormation template. CDK support for Amazon Managed Grafana is limited as it is a relatively new service offering, which has a heavy reliance on custom resources during deployment. Resource dependencies created outside the scope of the module are managed by Systems Manager parameters. These parameters are then used in IAM Identity Center policies to give appropriate permissions to other resources. The following steps are performed in the deployment:

EV Battery Health module 41

- **Create Amazon Managed Grafana workspace** An Amazon Managed Grafana workspace is created with the authentication provider set to AWS IAM Identity Center.
- Create Amazon Managed Grafana API key An Admin API key is created for the Amazon Managed Grafana workspace and is stored in Secrets Manager.
- Create Amazon Athena data source Amazon Athena is added as a data source to the workspace with the Amazon S3, AWS Glue, and Amazon Athena resources used for storing and indexing telemetry data taken as input to the module.
- Create dashboard An Amazon S3 bucket to store Amazon Managed Grafana assets is created. A custom resource creates the dashboard JSON data model and uploads it to the bucket. The PutObject action in the bucket invokes a Lambda function, which calls the Amazon Managed Grafana HTTP API to create the dashboard in the workspace.
- **Provision alerting** The Amazon Managed Grafana workspace does not have alerting capabilities by default. To support alerting, additional updates to the workspace configuration are made with custom resources post-deploy.
- Create alert rules A custom resource creates the alert rules JSON data model and uploads it to the Amazon Managed Grafana assets bucket. The PutObject action in the bucket triggers a Lambda function, which calls the Amazon Managed Grafana HTTP API to create the alert rules in the workspace.
- Create alert contact points and notification policy Alert rules must be configured with a
 contact point (where the alert messages are routed to) and a notification policy (which alert rules
 are routed to which contact points). The CMS EV Battery Health module uses Amazon SNS as the
 contact point for alerts.

EV Battery Health module 42

Vehicle Simulator module



The Vehicle Simulator module helps customers test vehicle communication with AWS IoT Core without the need for physical vehicles. This module provides a web-based GUI (Graphical user interface) that allows customers to create and simulate fleets of connected vehicles.

Simulated fleets emit data payloads at configured intervals from either a user-defined template or the default provided VSS schema. Simulated vehicles also onboard themselves with AWS IoT Core. You can monitor the simulation and observe how CMS on AWS services are processing the data.

Vehicle Simulator module 43

The GUI for the CMS Vehicle Simulator module is hosted as a static webpage in Amazon S3 and distributed through Amazon CloudFront. When deploying the module, an Amazon Cognito user pool is created with an initial user to enable authorization.

A REST API is implemented through API Gateway and backed by Lambda functions. This API provides functionality to create and store devices and simulations into DynamoDB tables.

Invoking the simulation API invokes the start of an AWS Step Functions workflow, which does the following:

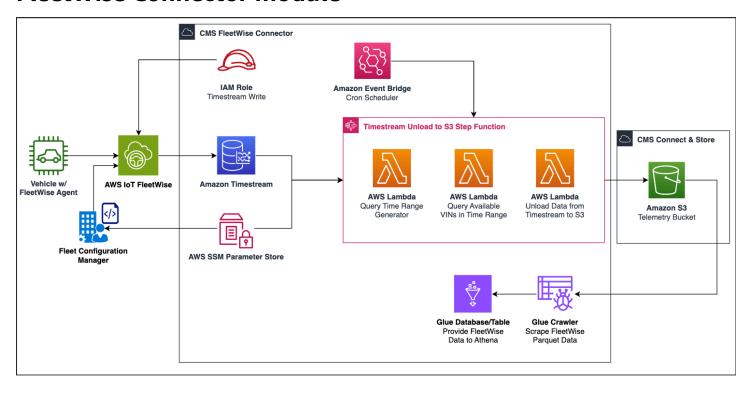
- 1. Maps device types to their corresponding payload configuration.
- 2. Retrieves payload configurations from a DynamoDB table.
- 3. Maps the total device count to Lambda function invocations.
- 4. For each Lambda function invocation, publishes the simulated payload to an AWS IoT MQTT topic.
- 5. Checks the overall duration and if elapsed, goes to step 7. If not elapsed, goes to step 6.
- 6. Waits the configured delay interval, then repeats steps 4 and 5.
- 7. Updates the DynamoDB table to represent the finished simulation state.

Important

This module is designed to simulate data for testing. It is not recommended for production environments.

Vehicle Simulator module

FleetWise Connector module



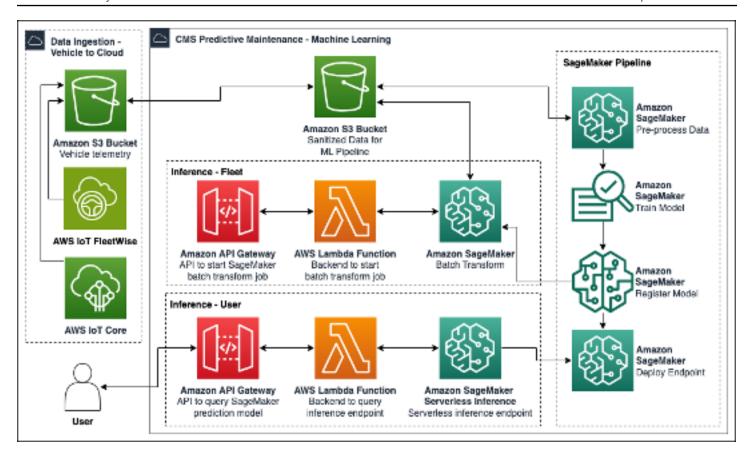
The FleetWise Connector module provides the required resources and roles to consume data that is captured by <u>AWS IoT FleetWise</u> campaigns. This is done by leveraging the Amazon Timestream's *Unload to S3* feature to migrate data to the CMS Connect & Store module's telemetry bucket on a recurring interval, orchestrated by AWS Step Functions. The data is then indexed using an AWS Glue Crawler job and made accessible through Amazon Athena.

In order for the CMS FleetWise Connector module to function, <u>AWS IoT FleetWise</u> must be configured and running. When configuring an AWS IoT FleetWise Campaign, the data store must be configured to the Timestream database, table, and role created by the CMS FleetWise Connector module during deploy. This Timestream database and table are used by default for the *Unload to S3* operation. SSM parameters are created with the Timestream database, table, and role that can be used in conjunction with automated scripting of AWS IoT FleetWise Campaign creation.

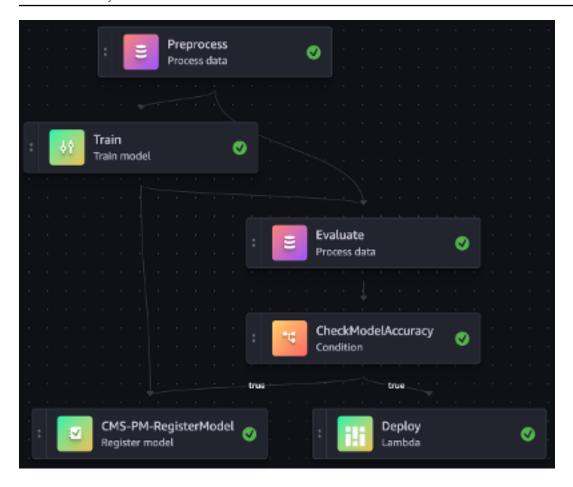
Predictive Maintenance module

The Predictive Maintenance module implements an MLOps framework to train, evaluate, retrain, and deploy neural-network based ML models for predicting vehicle faults. The following image shows the architecture of this module, which uses Amazon S3, Amazon SageMaker, Lambda, and API Gateway to create, store data for, execute pipelines on, and query these models.

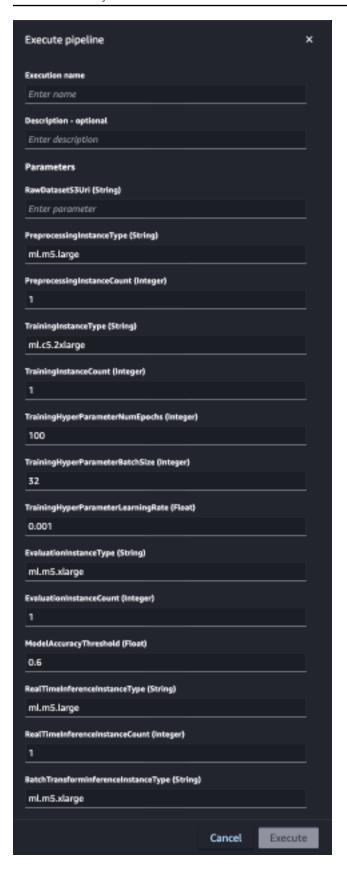
FleetWise Connector module 45



The MLOps framework is implemented using Amazon SageMaker Pipelines, as shown in the followings image. Amazon SageMaker Pipelines provides a mechanism to implement various stages in the pipeline. You can configure Amazon SageMaker Pipelines when creating a pipeline execution.

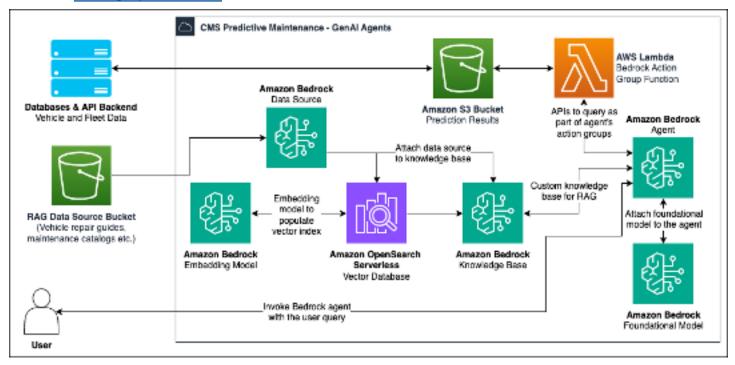


The following image shows the configuration options of the pipeline execution, after which the stages are described in more detail.



- Data The dataset for model training is expected in the form of a CSV file named dataset.csv. Every column except the last column is treated as model input, and the last column is treated as the output label. Upload the dataset to the S3 bucket created by the CMS Predictive Maintenance module. The name of the S3 bucket is stored in the Systems Manager parameter named /solution/cms/predictive-maintenance/predictor/pipeline/assets-bucket/name. Provide the dataset's Amazon S3 unique resource identifier (URI) as input to the pipeline in the RawDatasetS3Uri input parameter.
- **Pre-process** The pre-process stage takes the Amazon S3 dataset and splits it in three parts:
 - Training dataset Used in the training stage to train the model
 - Validation dataset Used during training stage to validate accuracy and avoid overfitting
 - Test dataset Used in the evaluation stage to evaluate the model's accuracy on data unseen during training
- Training In the training stage, a classification model is trained using PyTorch. The training
 dataset and the validation dataset created in the pre-process stage are both used in the
 training stage. The trained model is then stored and uploaded to Amazon S3 to be used in the
 subsequent stages of the pipeline.
- **Evaluate** In the evaluation stage, the trained model is used to perform inference by using the test dataset created in the pre-process stage. The model accuracy is then stored in an evaluation results file.
- Check model accuracy During the creation of the pipeline, you specify an input parameter
 called ModelAccuracyThreshold, which is the acceptable minimum threshold for model
 accuracy. The evaluation results stored during the evaluation stage are checked against the
 specified threshold. They pass the stage if the threshold value is met.
- **Register model** Register the trained model in a <u>Model Group</u>, which provides versioning for the models between multiple runs of the pipeline.
- **Deploy** In this stage, a Lambda function invokes the creation of a deployable Amazon SageMaker AI model and an endpoint that is used for real-time inference.
- APIs The module implements two APIs:
 - /predict API to provide the model with input data and get an output prediction.
 - /batch-predict to run batch inference on test data stored in the bucket with name stored
 in the Systems Manager parameter called /solution/cms/predictive-maintenance/
 predictor/pipeline/assets-bucket/name and its object key being latest.csv. The
 output of the batch prediction is stored in the same bucket with object key inference/
 latest.csv.out.

The APIs are authorized by a custom Lambda authorizer, which uses resources set up by the Auth Setup module and the CMS Auth module. The API requests are throttled with the default rate limit that is set at an account level. For implementing more sophisticated hardening against DDoS, consider setting up AWS WAF.



The Predictive Maintenance module provides a generative AI chatbot implemented using Amazon Bedrock Agents. With the chatbot, you can use natural language queries to get fleet insights from textual documents by using RAG. You can also use a function calling ability to guery backend databases and receive accurate and real-time vehicle related data. In addition, you can version the Amazon Bedrock Agent, and manage these versions with aliases. This can be done with the Amazon Bedrock console.

For RAG, the module uses Amazon Bedrock Knowledge Bases, which leverages an Amazon OpenSearch Service Serverless vector database to index textual documents stored in an S3 bucket.

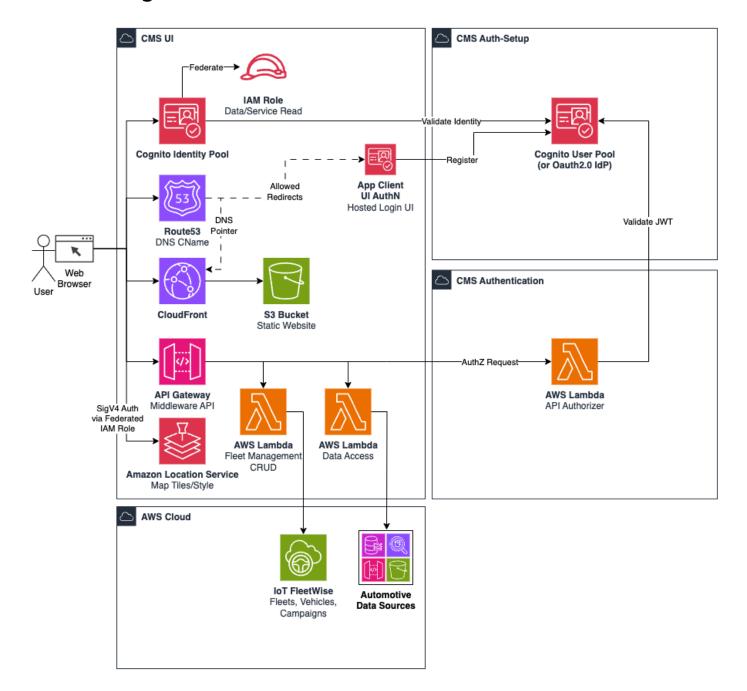
For function calling, the chatbot uses action groups that call Lambda functions to query the S3 bucket where the prediction results are stored to answer whether a vehicle with a given VIN requires maintenance.

Important

When using the generative AI chatbot in production environments, we recommend that you use Amazon Bedrock Guardrails to implement custom safeguards to control the outputs of

the chatbot according your compliance needs. The outputs of generative AI chatbots are not deterministic and can be susceptible to hallucinations.

Fleet Management UI module



Fleet Management UI module 51

The Fleet Management UI module integrates <u>AWS IoT FleetWise</u> to provide a web-based GUI (Graphical user interface) that allows management and monitoring of fleets and connected vehicles.

The GUI for the CMS Fleet Management UI module is hosted as a static webpage in Amazon S3 and distributed through Amazon CloudFront. When deploying the module, an Amazon Cognito user pool app client is created to enable authentication. Refer to the <u>Auth Setup</u> module for information on configuring a CMS compatible Cognito IdP. In addition, a Cognito Identity Pool is created for authentication to Amazon Location Services for use with the embedded map views.

A REST API is implemented through API Gateway and backed by Lambda functions. This API provides functionality to:

- 1. View and manage fleets.
- 2. View and manage vehicles.
- 3. View and manage fleet and vehicle associations.
- 4. View campaigns.
- 5. Start, stop, and delete campaigns associated with a fleet.

The UI also contains some concept views and widgets which are not currently connected to live data. These are included as samples that showcase how collected vehicle data can be utilized to build fleet monitoring functionality.

The sample views and widgets are:

- Fleet dashboard a collection of sample widgets showcasing ways connected vehicle data could be utilized to monitor a fleet.
 - Fleet driver scores
 - Vehicle health status
 - Vehicle utilization status
 - Miles driven per day
 - · Battery state of health
 - Hard breaking events
 - Vehicle ADAS version
- Map View a sample map view powered by Amazon Location Services. Interactive vehicles move along preset routes on the map. Vehicle routes and events can be seen by clicking on them.

Fleet Management UI module 52

- Vehicle dashboard a collection of sample widgets showcasing ways connected vehicle data could be utilized to monitor vehicle state and health.
 - Vehicle status
 - Vehicle location
 - Operating hours
 - · Battery state of health
 - · Vehicle speed
 - Vehicle alerts
- Maintenance Alerts sample alerts that could be collected from vehicles in a fleet.



This module is currently in a preview state. It is not recommended for production environments.

AWS services in this solution

| AWS service | Description |
|-------------------------------|--|
| Amazon API Gateway | Core. Hosts REST API endpoints in the solution. |
| AWS AppSync | Core. Provides GraphQL API endpoints in the solution. |
| Amazon Athena | Core. Performs queries on vehicle data stored in Amazon S3. |
| AWS Certificate Manager (ACM) | Core. Generates certificates used for validation of HTTPS requests. |
| AWS Chalice | Core. Framework used to define and deploy the Vehicle Simulator module's API. |

| AWS service | Description |
|-------------------------------------|---|
| AWS Cloud Development Kit (AWS CDK) | Core. Enables infrastructure as code for the entirety of CMS on AWS. |
| AWS CloudFormation | Core. Manages deployments for the solution infrastructure. |
| Amazon CloudWatch | Core. The solution's code and infrastructure emit logs to Amazon CloudWatch. |
| AWS CodeBuild | Core. Defines build steps to aid in deploying and managing the solution. |
| AWS CodePipeline | Core. Runs collections of AWS CodeBuild steps for deploying and managing the solution. |
| Amazon Data Firehose | Core. High throughput writing and transform ation of MQTT topic payloads to Amazon S3. |
| Amazon DynamoDB | Core. Primary non-relational data storage used for vehicle and user records. |
| AWS Glue | Core. Validates and transforms JSON formatted payloads of a specified schema into Parquet format. |
| <u>AWS IAM</u> | Core. Authorizes CMS on AWS resources and users throughout solution with associated least-privilege roles and policies. |
| AWS IoT Core | Core. Primary service for onboarding and connecting vehicles, as well as management of those vehicles (things), their certificates, and policies. Also used for data ingestion through MQTT topics for communication between modules. |

| AWS service | Description |
|---|---|
| AWS KMS | Core. Encrypts data in transit and at rest throughout the solution. |
| AWS Lambda | Core. A variety of runtime logic with serverless code throughout the solution. |
| Amazon Managed Grafana | Core. Build dashboards and configure alerts based on vehicle data stored in Amazon S3. |
| Parameter Store, a capability of AWS Systems Manager | Core. Stores configuration level information used throughout the solution. |
| Amazon S3 | Core. General purpose shared data storage used throughout CMS on AWS. |
| AWS Secrets Manager | Core. Stores and rotates secrets used throughout the solution. |
| Amazon SNS | Core. Used to publish and subscribe to messages. |
| Amazon SQS | Core. Used to deliver messages between modules. |
| AWS Step Functions | Core. Orchestrates and manages Lambda functions throughout solution. |
| Amazon Timestream | Core. Timestream database used within the FleetWise Connector module. |
| Amazon VPC | Core. Networking boundary used throughout the solution. |
| Amazon Aurora PostgreSQL-compatible | Supporting. Database used to aid with deployment. |
| Amazon Bedrock | Supporting. Used by the Predictive Maintenan ce module. |

| AWS service | Description |
|--------------------------------------|---|
| Amazon CloudFront | Supporting. Provides a domain name to serve static web content and reduce latency of API endpoints. |
| Amazon ECS | Supporting. Simplifies the deployment, management, and scaling of the Backstage module. |
| Amazon ECR | Supporting. Image repository for Dockerized containers used by the Backstage module. |
| Elastic Load Balancing | Supporting. Provides network connections to Amazon ECS tasks running Backstage to aid with deployment. |
| Amazon EventBridge | Supporting. Used within the FleetWise Connector module for scheduled execution of the Step Functions flow. |
| AWS Fargate | Supporting. Combined with ECS, aids in management of EC2 instances within the Backstage module. |
| AWS IAM Identity Center | Supporting. Service that provides authentic ation for the Amazon Managed Grafana workspace. |
| AWS IoT FleetWise | Supporting. Used by the FleetWise Connector module. |
| Amazon OpenSearch Service Serverkess | Supporting. Used by the Predictive Maintenan ce module. |
| Amazon Route 53 | Supporting. Domain hosting integration used by the Backstage module. |

| AWS service | Description |
|-------------------------|--|
| Amazon SageMaker AI | Supporting. Used by the Predictive Maintenan ce module. |
| AWS X-Ray | Supporting. Traces runs of the AWS Step Functions. |
| Amazon Cognito | Optional. Authenticates users and internal services across the solution. |
| Amazon Location Service | Optional. Geographical functionality for the Vehicle Simulator and Fleet Management UI modules. |

Plan your deployment

This section describes the Regions, <u>cost</u>, <u>security</u>, and <u>quota</u> considerations before deploying the solution.

Supported AWS Regions

For the most current availability of AWS services by Region, see the AWS Regional Services List.

Connected Mobility Solution on AWS is supported in the following AWS Regions:

| Region name | |
|-----------------------|-----------------------|
| US East (Ohio) | Asia Pacific (Tokyo) |
| US East (N. Virginia) | Asia Pacific (Sydney) |
| US West (Oregon) | Europe (Frankfurt) |
| Europe (Ireland) | |

Cost

You are responsible for the cost of the AWS services used while running this solution. As of this revision, the cost for running this solution with the default settings in the US East (N. Virginia) is approximately \$352.80 a month.

See the pricing webpage for each AWS service used in this solution.

We recommend creating a <u>budget</u> through AWS Cost Explorer to help manage costs. Prices are subject to change. For full details, see the pricing webpage for each <u>AWS service used in this solution</u>.

CMS on AWS static cost tables

CMS on AWS total

Supported AWS Regions 58

| AWS service | Amount | Dimensions | Cost [USD] |
|---|--------|----------------------------------|------------|
| Amazon VPC | 17 | Endpoints in 2 AZs | |
| | 2 | NAT gateways | \$314.00 |
| AWS KMS | 29 | Customer managed keys | \$29.00 |
| AWS Secrets Manager | 7 | Sensitive data stored in secrets | \$2.80 |
| Amazon Managed Grafana | 1 | 1 Workspace × API license | \$9.00 |
| Amazon OpenSearch Service Serverless | 1 | 1 Collection (\$0.48 per hour) | \$345.20 |

Config module

VPC module

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|--------------------|------------|
| Amazon VPC | 17 | Endpoints in 2 AZs | |
| | 2 | NAT gateways | \$314.00 |

Auth Setup module

| AWS service | Amount | Dimensions | Cost [USD] |
|---------------------|--------|-------------------|------------|
| AWS Secrets Manager | 3 | PostgreSQL secret | \$1.20 |

CMS on AWS static cost tables 59

Automotive Cloud Developer Portal

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|--------------------------|------------|
| AWS KMS | 5 | Customer managed keys | \$5.00 |

• Backstage module*

| AWS service | Amount | Dimensions | Cost [USD] |
|---------------------|--------|--------------------------|------------|
| AWS KMS | 1 | Customer managed keys | \$1.00 |
| AWS Secrets Manager | 2 | PostgreSQL secret | \$0.80 |

• Auth module*

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|------------|------------|
|-------------|--------|------------|------------|

Vehicle Provisioning module

| AWS service | Amount | Dimensions | Cost [USD] |
|---------------------|--------|-------------------------------------|------------|
| AWS KMS | 2 | Customer managed keys | \$2.00 |
| AWS Secrets Manager | 1 | Claim certificate created on deploy | \$0.40 |

Connect and Store module

CMS on AWS static cost tables 60

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|--------------------------|------------|
| AWS KMS | 3 | Customer managed keys | \$3.00 |

API module

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|-----------------------|------------|
| AWS KMS | 1 | Customer managed keys | \$1.00 |

Alerts module

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|--------------------------|------------|
| AWS KMS | 7 | Customer managed keys | \$7.00 |

• EV Battery Health module*

| AWS service | Amount | Dimensions | Cost [USD] |
|---------------------|--------|---|------------|
| AWS KMS | 3 | Customer managed keys | \$3.00 |
| AWS Secrets Manager | 1 | Amazon Managed Grafana API key secret | \$0.40 |
| AWS Grafana | 1 | 1 Workspace x API license | \$9.00 |

• Vehicle Simulator module*

CMS on AWS static cost tables 61

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|-----------------------|------------|
| AWS KMS | 2 | Customer managed keys | \$2.00 |

FleetWise Connector module

| AWS service | Amount | Dimensions | Cost [USD] |
|-------------|--------|-----------------------|------------|
| AWS KMS | 3 | Customer managed keys | \$3.00 |

Predictive Maintenance module

| AWS service | Amount | Dimensions | Cost [USD] |
|---|--------|--------------------------------|------------|
| AWS KMS | 2 | Customer managed keys | \$2.00 |
| Amazon OpenSearch Service Serverless | 1 | 1 Collection (\$0.48 per hour) | \$345.20 |

Example costs

This solution provides zero-cost-on-no-use. Consider the following use cases:

Use case 1 - Solution deployed for 100 devices connected with little to no analytics run.

| Service | Configuration summary | Monthly Cost [USD] |
|---------|--|--------------------|
| MQTT | Number of devices (MQTT) (100), Average size of each message (5 KB), Average size of each record (1 KB), Average size of each record | \$2.59 |

| Service | Configuration summary | Monthly Cost [USD] |
|-------------|---|--------------------|
| | (1 KB), Average number of actions applied per rule (2), Average size of each message (5 KB), Average number of connection minutes for a device (24000), Number of messages for a device (24000), Number of rules initiated (5). | |
| S3 Standard | S3 Standard storage (1 GB per month), PUT, COPY, POST, LIST requests to S3 Standard (720), GET, SELECT, and all other requests from S3 Standard (720), Data returned by S3 Select (1 GB per month), Data scanned by S3 Select (4 GB per month). | \$0.04 |

| Service | Configuration summary | Monthly Cost [USD] |
|-----------------------------|--|--------------------|
| Amazon Data Firehose | Source Type (Direct PUT or Kinesis Data Stream), Data records units (exact number), Record size (5 KB), Dynamic Partitioning (Add On) (Enabled), Data format conversion (optional) (Enabled), Average ratio of data processed to VPC vs data ingested (1.3), Number of records for data ingestion (1200 per day), Number of subnets for VPC delivery (0), Average size objects delivered (64 MB), JQ Processing (optional) (Enabled), Average JQ expected processing hours (70). | \$4.91 |
| Amazon Athena | Amount of data scanned per query (1 GB), Total number of queries (100 per day) | \$14.85 |
| DynamoDB on-demand capacity | Table class (Standard), Average item size (all attributes) (3 KB), Data storage size (0.1 GB) | \$0.05 |
| Amazon CloudFront | Data transfer out to internet (1 GB per month), Number of requests (HTTPS) (1000 per month). | \$0.09 |

| Service | Configuration summary | Monthly Cost [USD] |
|-------------------------|---|--------------------|
| Amazon CloudWatch | Standard Logs: Data Ingested (0.1 GB), Logs Delivered to CloudWatch Logs: Data Ingested (0.1 GB), Logs Delivered to Amazon S3: Data Ingested (0.1 GB), Number of Lambda functions (20), Number of requests per function (10 per day). | \$48.13 |
| Amazon Cognito | Number of monthly active users (MAU) (100), Advanced security features (Enabled) | \$5.00 |
| Amazon Managed Grafana | Number of active editors/ administrators (1 per workspace per month), Number of active viewers (5 per workspace per month). | \$34.00 |
| Standard topics | DT Inbound: Not selected (0 TB per month), DT Outbound: Not selected (0 TB per month), Requests (10000 per month), EMAIL/EMAIL-JSON Notifications (10000 per month), Mobile Push Notificat ions (20000 per month). | \$0.19 |
| AWS AppSync API Request | Number of API Requests (10 thousand per month). | \$0.04 |

Use case 2 - Solution deployed with the following parameters.

1. Cost estimated for 1 million vehicles.

- 2. Certificates are rotated annually. Consider exploring non-AWS options to host your own Certificate Authority (CA).
- 3. A baseline of 10 alerts per vehicle per day.
- 4. Data ingestion to the cloud, approximately 10MB per vehicle per day.
- 5. 50 signals per second (2,000) per vehicle for 5x20 hours a month 5KB of data per minute; maintaining the message size at or below 5KB is recommended to save costs.
- 6. Monitoring of 100 Diagnostic Trouble Codes (DTC) per vehicle per month for vehicle health monitoring.
- 7. 10-15 trips per vehicle per day, with 1MB of vehicle data stored in DynamoDB storage. Historical data will be exported to Amazon S3. Vehicle profile data will reside in DynamoDB, while all vehicle-generated data will be stored in Amazon S3.
- 8. 100 Simple Notification Service (Amazon SNS) notification requests per vehicle per day, alongside 20 mobile push notifications per vehicle per day.
- 9. Five Dashboard editors and 50 viewers for Grafana.
- 10In the context of DevSecOps, we anticipate retaining logs for a period of three months. Please be aware that the cost of CloudWatch will increase because we are obligated to retain logs for more than one month, with a minimum retention period of three months.
- 11. The cost of Amazon S3 includes the monthly cost of Standard S3, which amounts to \$4,000.00 every month.

| AWS service | Dimensions | Monthly Cost [USD] |
|-------------|--|--------------------|
| MQTT | Number of devices (MQTT) (1200000), Average size of each message (5 KB), Average size of each record (1 KB), Average size of each record (1 KB), Average number of actions taken per rule (2), Average size of each message (5 KB), Average number of connection minutes for a device (28000), Number of messages for a device | \$26,910.7 |

| AWS service | Dimensions | Monthly Cost [USD] |
|----------------------|---|--------------------|
| | (28000), Number of rules triggered (5) | |
| S3 Standard | S3 Standard storage (10 TB per month), PUT, COPY, POST, LIST requests to S3 Standard (920000000), GET, SELECT, and all other requests from S3 Standard (920000000), Data returned by S3 Select (5 TB per month), Data scanned by S3 Select (8 TB per month) | \$5,249.09 |
| Amazon Data Firehose | Source Type (Direct PUT or Kinesis Data Stream), Data records units (exact number), Record size (5 KB), Dynamic Partitioning (Add On) (Enabled), Data format conversion (optional) (Enabled), Average ratio of data processed to VPC vs data ingested (1.3), Number of records for data ingestion (1200000000 per day), Number of subnets for VPC delivery (0), Average size objects delivered (64 MB), JQ Processing (optional) (Enabled), Average JQ expected processing hours (70) | \$11,679.87 |

Example costs 67

| AWS service | Dimensions | Monthly Cost [USD] |
|-----------------------------|--|--------------------|
| Amazon Athena | Amount of data scanned per query (100 GB), Total number of queries (110000 per month) | \$53,710.94 |
| DynamoDB on-demand capacity | Table class (Standard), Average item size (all attributes) (3 KB), Data storage size (10 GB) | \$185.00 |
| Amazon CloudFront | Data transfer out to internet (1 TB per month), Number of requests (HTTPS) (60000 per month), Data transfer out to origin (1 TB per month) | \$107.58 |
| Amazon CloudWatch | Standard Logs: Data Ingested (1 TB), Logs Delivered to CloudWatch Logs: Data Ingested (10 GB), Logs Delivered to S3: Data Ingested (10 GB), Number of Lambda functions (20), Number of requests per function (10000 per day) | \$575.64 |
| Amazon Cognito | Number of monthly active users (MAU) (100), Advanced security features (Enabled) | \$5.00 |
| Amazon Managed Grafana | Number of active editors/ administrators (5 per workspace per month), Number of active viewers (50 per workspace per month) | \$295.00 |

Example costs 68

| AWS service | Dimensions | Monthly Cost [USD] |
|--------------------------------------|---|--------------------|
| Standard topics | Requests (200 million per month), EMAIL/EMAIL-JSON Notifications (200 million per month), Mobile Push Notificat ions (20 million per month) | \$4,109.48 |
| AWS AppSync API Request | Number of API Requests (300 million per month) | \$1,200.00 |
| Amazon Elastic Container Registry | DT Inbound: All other regions (1 TB per month), DT Outbound: Not selected (0 TB per month), Amount of data stored (280 GB per month) | \$28.00 |

Note

The price is based on public pricing as displayed in the AWS Price Calculator. Customers may have a special pricing agreement based on reserve capacity and negotiated rates with AWS.

Security

When you build systems on AWS infrastructure, security responsibilities are shared between you and AWS. This shared responsibility model reduces your operational burden because AWS operates, manages, and controls the components including the host operating system, the virtualization layer, and the physical security of the facilities in which the services operate. For more information about AWS security, visit AWS Cloud Security.

As part of this shared responsibility model, we implement the security best practices of the AWS Well-Architected Framework. More details on these security considerations can be found in the Security section of the Architecture overview. Some services have further security considerations available, which are purposefully omitted to allow for customization and configuration based on your security needs. See the following subsections for more detail.

Security

Important

This solution is not designed to handle personally identifiable information (PII). See PII data for more information.

Authentication and authorization

All CMS on AWS API operations are protected through authentication requirements for both users and services. Both users and services must provide a valid access token associated with the IdP configured within the solution. By allowing customers to use their own identity provider, they have full control over the configuration of their authentication system.

All JWTs used for authentication and authorization are validated through protocol defined by OAuth2.0 standards. For more details, see Auth module.

The authentication flow is protected against security risks and attacks by implementing a variety of safety procedures. These include the use of client secrets for both the user and service app client, an optional PKCE code verifier for user authentication, and the use of the authorization code flow for user authentication.

Amazon API Gateway

This solution deploys an Amazon API Gateway REST API and uses the default API endpoint and SSL certificate. The default API endpoint supports TLSv1 security policy. To use a later version of TLS, use your own domain name and custom SSL certificate. For more information, refer to Choosing a minimum TLS version for a custom domain in API Gateway in the Amazon API Gateway Developer Guide.

Amazon Bedrock

This solution's default configuration doesn't deploy Amazon Bedrock Guardrails for the Amazon Bedrock Agents. This is largely due to the modular nature of our solution, and the ease at which you can opt out of deploying the Predictive Maintenance module.

We recommend that you enhance your generative AI security by setting up Amazon Bedrock Guardrails manually. Amazon Bedrock Guardrails provides additional customizable safeguards on top of the native protections of FMs, and allow your customers to opt-out of generative AI use.

Authentication and authorization 70

71

Content Security Policy (CSP)

Content Security Policy (CSP) is a security standard that helps protect web applications by controlling the sources from which content like scripts, images, and styles can be loaded. It reduces the risk of attacks like Cross-Site Scripting (XSS) by allowing website administrators to specify which domains are trusted. While the default CSP provides a robust layer of security, we recommend customizing the CSP configuration of your deployment based on your specific security requirements. Tailoring the policy to fit the unique structure and needs of your web applications ensures a more effective defense against evolving threats.

Amazon CloudFront

This solution deploys a web console <u>hosted</u> in an Amazon S3 bucket. To help reduce latency and improve security, this solution includes a CloudFront distribution with an origin access identity, which is a CloudFront user that provides public access to the solution's website bucket contents. For more information, see <u>Restricting access to an Amazon S3 origin</u> in the *Amazon CloudFront Developer Guide*.

Amazon CloudFront is deployed using the default CloudFront domain name and TLS certificate. The default CloudFront SSL certificate supports TLSv1. To use a later TLS version, use your own custom domain name and custom SSL certificate. For more information, refer to <u>Using alternate</u> domain names and HTTPS in the *Amazon CloudFront Developer Guide*.

Amazon CloudWatch

You can use CloudWatch to set up important security features such as <u>canaries</u> and <u>alarms</u>. CloudWatch is not configured by default in this solution. However, modules are registered in the AWS Service Catalog AppRegistry and viewable with the myApplications dashboard. This dashboard includes a widget for CloudWatch, where you can create, manage, and view your canaries and alarms.

We recommend configuring CloudWatch to monitor and protect your solution. For more details, see <u>Using Amazon CloudWatch alarms</u>.

Customer managed AWS KMS keys

This solution uses encryption at rest for securing data and employs <u>customer managed keys</u> for customer data and AWS managed keys for AWS service data. These keys are used to automatically and transparently encrypt your data before it is written to storage layers. Some users might prefer

Content Security Policy (CSP)

to have more control over their data encryption processes. This approach allows you to administer your own security credentials, offering a greater level of control and visibility.

Amazon DynamoDB

This solution deploys DynamoDB tables to your account without point-in-time recovery (PITR) enabled by default. PITR allows you to restore your DynamoDB tables to any moment in time within the last 35 days. It provides an automatic, continuous backup of your data, offering protection against accidental writes or deletes. PITR is disabled by default to protect against the exposing of sensitive customer data upon teardown. We recommend enabling PITR within your deployment to protect your data against accidental loss.

Elastic Load Balancing

This solution deploys Elastic Load Balancing. By default, deletion protection is not enabled on Elastic Load Balancing. This protects against accidental resource retention in your AWS accounts and assists with teardown. For production deployments, we recommend enabling deletion protection and using <u>AWS Config</u> to monitor and detect changes in the deletion protection status of Elastic Load Balancing.

PII data

This solution is not designed with the advanced security protocols necessary to store, process, or handle PII. All data is encrypted in-transit and at rest; however, this solution doesn't vet or filter incoming data for PII elements. As a result, you must ensure that no PII is included in the data transmitted.

Amazon SageMaker Al

This solution deploys Amazon SageMaker AI. Consider following <u>Amazon SageMaker AI security</u> <u>best practice guidance</u> for data privacy, data protection, logging, and monitoring considerations. We also recommend considering enabling AWS Config to surface common Amazon SageMaker AI security misconfigurations within your account.

AWS WAF

This solution's default configuration doesn't deploy a web application firewall (WAF) in front of the API endpoints. To enhance your API security by setting up a WAF, you must do so manually. AWS

Amazon DynamoDB 72

provides an in-depth guide on how you can control access to your API Gateway with AWS WAF. For instructions on how to implement AWS WAF in front of your API and increase distributed denial of service (DDoS) protection for your web applications, see Using AWS WAF to protect your APIs.

Other AWS services

Individual AWS services have additional security best practices and considerations, most of which are configured by default in this solution. The following table provides details of these security considerations. You can find further details about the usage of each service in the <u>AWS services in this solution</u> section. We recommend reviewing the security considerations and details for services that are relevant to your use case and security needs.

| AWS service | Security considerations for this solution |
|---|---|
| Amazon API Gateway | Ensure that API Gateway authorization settings are properly configured (for example, IAM, Lambda authorizers). Use throttling to mitigate DDoS attacks. |
| AWS AppSync | AWS AppSync requires secure authentication (for example, API keys, IAM, Amazon Cognito). Be cautious of GraphQL queries, which can expose large datasets if not managed properly. |
| Amazon Athena | Protect sensitive data queries by using encryption. Ensure proper IAM permissions for access to data in Amazon S3. |
| Amazon Aurora PostgreSQL-Compatible Edition | Use AWS KMS for encryption at rest. Configure network isolation with Amazon VPC. Apply least privilege IAM roles for accessing database instances. |
| Amazon Bedrock | Restrict AI model access by using IAM policies. Consider data privacy issues when using generative AI. Secure communication with TLS. |

| AWS service | Security considerations for this solution |
|-------------------------------|--|
| AWS Certificate Manager (ACM) | Use ACM to manage SSL/TLS certificates, ensuring data in transit is encrypted. Regularly rotate certificates and manage access to certificate requests with fine-grained IAM policies. |
| AWS Chalice | Ensure that permissions in the deployment policy are minimal and follow least privilege practices. Use IAM roles for AWS resource access. |
| AWS CDK | Apply least privilege IAM roles. Encrypt sensitive information in environment variables or configurations. |
| AWS CloudFormation | Carefully manage IAM roles and permissions in templates to ensure least privilege. Control access to stacks and avoid including sensitive information in CloudFormation templates. |
| Amazon CloudFront | Use HTTPS for secure content delivery. Apply IAM roles for access control. Configure origin access control to prevent direct access to your S3 bucket. |
| AWS CloudTrail | Enable CloudTrail logging to help you with governance, compliance, operational auditing, and auditing of your AWS account. |
| Amazon CloudWatch | Encrypt logs and metrics at rest. Use fine- grained IAM policies to control access to log groups and avoid exposing sensitive data in logs. |

| AWS service | Security considerations for this solution |
|------------------------|---|
| AWS CodeBuild | Limit access to the build environment and apply IAM roles with least privilege. Use encryption for build artifacts and ensure that sensitive data is not exposed in build logs. |
| AWS CodePipeline | Ensure least privilege for pipeline access and encrypt sensitive artifacts. Use secure endpoints (HTTPS) for integration with third- party tools. |
| AWS Config | Use MFA for user sign-ins. Control access to user pools by using IAM roles. Encrypt user data. |
| Amazon Data Firehose | Encrypt data both in transit and at rest. Configure IAM policies to control access to delivery streams. Validate that sensitive data is handled appropriately. |
| Amazon DynamoDB | Implement IAM policies to control access. Enable encryption at rest and in transit. Configure fine-grained access control to limit data exposure. |
| Amazon ECS | Use IAM roles for Amazon ECS tasks and services. Secure communication between tasks by using security groups. Encrypt sensitive data at rest. |
| Amazon ECR | Control access to Amazon ECR repositories by using IAM policies. Encrypt images in transit and at rest. Scan images for vulnerabilities. |
| Elastic Load Balancing | Use HTTPS for secure communication. Configure security groups to limit traffic. Enable logging for traffic monitoring. |

| AWS service | Security considerations for this solution |
|---------------------|--|
| Amazon EventBridge | Control access to event buses by using IAM policies. Encrypt sensitive data in event payloads. Use VPC endpoints for secure private communication. |
| AWS Fargate | Use IAM roles to control access to tasks. Apply network isolation with Amazon VPC. Encrypt data stored in Fargate tasks. Monitor task activity and limit container privileges. |
| AWS Glue | Configure fine-grained permissions for data sources and destinations. Enable encryption for data in transit and at rest. |
| IAM | Apply the principle of least privilege to all users and roles. Regularly rotate access keys. Enable MFA for critical accounts. |
| IAM Identity Center | Enforce MFA. Control permissions with least privilege policies. Regularly audit access to ensure compliance with security policies. |
| AWS IoT Core | Encrypt MQTT communication with TLS. Ensure that device policies follow least privilege. Use <u>AWS IoT Device Defender</u> for security audits and monitoring. |
| AWS IoT FleetWise | Implement device authentication and encryption for data in transit. Restrict access to fleet data by using IAM roles and policies. |
| AWS KMS | Control access to keys with fine-grained IAM policies. Enable logging for key usage and management actions. |

| AWS service | Security considerations for this solution |
|--|---|
| Lambda | Ensure that IAM policies are minimal and limit function execution access. Encrypt environme nt variables and monitor Lambda functions for unusual behavior. |
| Amazon Location Service | Encrypt geospatial data in transit and at rest. Control access to location data by using IAM roles. Review audit logs regularly for unusual access. |
| Amazon Managed Grafana | Implement secure user authentication. Integrate with IAM Identity Center or IAM roles. Monitor dashboards for data security and access control. Encrypt connections for data sources. |
| Amazon OpenSearch Service Serverless | Enable encryption for data at rest and in transit. Control access with fine-grained IAM policies. Monitor for unusual search activity. |
| Parameter Store, a feature of Systems Manager | Restrict access to sensitive parameters by using IAM policies. Enable encryption for secure string parameters with AWS KMS. Monitor parameter access through logging. |
| Amazon Route 53 | Encrypt DNS queries. Monitor for DNS hijacking attempts. Use IAM roles for managing hosted zones and access to DNS records. |
| Amazon S3 | Apply least privilege access policies to S3 buckets. Encrypt data at rest. Enable versionin g and logging for audit purposes. |

| AWS service | Security considerations for this solution |
|-------------------------------------|--|
| Amazon SageMaker AI | Use encryption for training data and models. Control access to resources by using IAM roles. Monitor access to SageMaker AI notebooks and endpoints. |
| AWS Secrets Manager | Restrict access to secrets by using IAM policies. Rotate secrets regularly. Encrypt secrets with AWS KMS. |
| Amazon SNS | Use IAM policies to restrict access to topics. Enable encryption for messages at rest. Use VPC endpoints for secure communication within a VPC. |
| Amazon SQS | Apply fine-grained access control using IAM policies. Encrypt messages with SSE. Configure VPC endpoints to restrict Amazon SQS access within a VPC. |
| AWS Step Functions | Implement appropriate IAM permissions for task execution. Avoid embedding sensitive data in state machines. Encrypt data at rest. |
| Amazon Timestream for LiveAnalytics | Configure fine-grained access control using IAM roles. Enable encryption at rest for data stored in Timestream, and secure data in transit using TLS. |
| Amazon VPC | Control access to resources within VPCs using security groups and NACLs. Use VPC Flow Logs for monitoring and ensure encryption for data in transit using TLS. |
| AWS X-Ray | Control access to traces by using fine-grained IAM policies. Ensure that sensitive data is not included in traces. Encrypt data at rest. |

Quotas

Service quotas, also referred to as limits, are the maximum number of service resources or operations for your AWS account.

Quotas for AWS Services in this solution

Make sure you have sufficient quota for each of the <u>services implemented in this solution</u>. For more information, see AWS service quotas.

Use the following links to go to the page for that service. To view the Service Quotas for all AWS services in the documentation without switching pages, view the information in the <u>Service</u> endpoints and quotas page in the PDF instead.

AWS CloudFormation quotas

Your AWS account has CloudFormation quotas that you should be aware of when deploying this solution. By understanding these quotas, you can avoid limitation errors that would prevent you from deploying this solution successfully. For more information, see AWS CloudFormation User's Guide.

Limitations

Some aspects of the solution have been optimized for initial cost management. Consider the following when using the solution in production:

- High availability ACDP/Backstage is designed to be used primarily as a development tool and therefore does not use a multi-AZ RDS instance. Consider enabling multiple instances of your database if you need high uptime guarantees.
- Backup and restore To configure backup and restore support for recovery scenarios, AWS
 Backup must be manually enabled after deploying ACDP/Backstage and CMS modules for each
 of the resources that require backup and restore functionality.

Quotas 79

Deploy the solution

This solution uses AWS CloudFormation templates and stacks to automate its deployment. The CloudFormation templates specify the AWS resources included in this solution and their properties. The CloudFormation stack provisions the resources that are described in the templates.

Deployment process overview

To directly deploy ACDP with Backstage and CMS on AWS modules into your AWS account, see the AWS CloudFormation templates section below to view, download, or deploy each CloudFormation template. Some parts of the solution have dependencies; therefore, a deployment order is roughly enforced. See the CMS module deployment order section below for details.

For more detailed instructions on deploying CMS on AWS modules from the solution's repository, and for instructions that might be better suited for developers, see the CMS on AWS README Deploy section.

For a detailed walkthrough of how to deploy CMS on AWS modules from ACDP, and more specifically Backstage, see the CMS on AWS README Deploy CMS Modules via Backstage section.

Before you launch the solution, review the cost, architecture, network security, and other considerations discussed earlier in this guide.

Time to deploy ACDP: Approximately 30–45 minutes

Time to deploy CMS on AWS modules: Varies based on module choice

Important

This solution includes an option to send anonymized operational metrics to AWS. We use this data to better understand how customers use this solution and related services and products. AWS owns the data gathered through this survey. Data collection is subject to the **AWS Privacy Notice.**

To opt out of this feature, download the template, modify the AWS CloudFormation mapping section, and then use the AWS CloudFormation console to upload your updated template and deploy the solution. For more information, see the Anonymized data collection section of this guide.

Deployment process overview

CMS module deployment order

All CMS on AWS modules have dependencies on the initial three deployments for configuring CMS on AWS. Some CMS on AWS modules have secondary dependencies on other modules and must be deployed in order. The rest of the modules do not have dependencies on other modules and can be deployed in any order after CMS Config.

Adhere to the following deployment order:

Deployment order of required CMS config

- 1. VPC
- 2. Auth Setup
- 3. CMS Config



Note

At this point, ACDP and Backstage should be deployed to assist with the further module deployments.

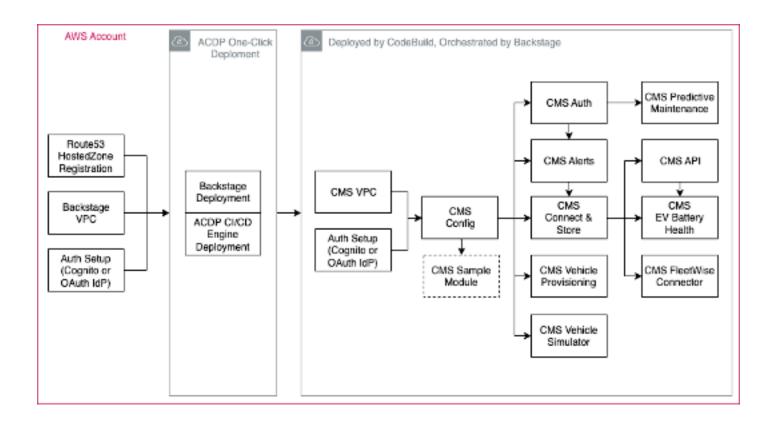
Deployment order of modules with dependencies

- 1. Auth
- 2. Alerts
- 3. Connect and Store
- 4. API
- 5. EV Battery Health
- 6. FleetWise Connector
- 7. Predictive Maintenance

Modules without dependencies

- Vehicle Provisioning
- Vehicle Simulator
- Fleet Management UI

This order is represented in the following diagram:



AWS CloudFormation templates

You can download the CloudFormation templates for this solution before deploying them.

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|-------------|----------------------|
| vpc.templ ate - Use this template to deploy the VPC module and all associate | view template button | launch button | None | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|-------------|----------------------|
| d component s. The default configuration deploys the services found in the VPC module section, and you can customize the template to meet your specific needs. | | | | |
| auth-setu p.templat e - Use this template to deploy the Auth Setup module and all associate d component s. The default configuration deploys the services found in the Auth Setup module section, and you can customize the template to meet your specific needs. | view template button | launch button | None | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-----------------------|---|---|--|
| acdp.temp late - Use this template to deploy the ACDP module and all associate d component s. The default configuration deploys the services found in the Automative Cloud Developer Portal section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc | Multi Account/M ulti Region Deployment 1. To Enable this feature, provide valid values to three CloudForm ation parameter s when you deploy this CloudForm ation Template. a. ENABLE_MU LTI_ACCOU NT_DEPLOY MENT: This is a boolean variable, use value true to enable the feature and |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---------------|-----------------------|---|-------------|---|
| | | | | false to disable. b. ORGS_MANA GEMENT_AW S_ACCOUNT _ID: This is the AWS Account from where you manage your Organizat ion or where the Organizat ion root is located. c. ORGS_MANA GEMENT_AC COUNT_REG ION: This is the region where |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---------------|-----------------------|---|-------------|---|
| | | | | you would deploy the which is explained later. 2. Upon successfull deploymen t follow the guidance given in the GitHub repository. |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|---|----------------------|
| cms-confi g.templat e - Use this template to deploy the Config module and all associate d component s. The default configuration deploys the services found in the Config module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|---|----------------------|
| cms-auth. template - Use this template to deploy the Auth module and all associate d component s. The default configuration deploys the services found in the Auth module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setupvpccms-config | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|---|----------------------|
| cms-alert s.template - Use this template to deploy the Alerts module and all associate d component s. The default configuration deploys the services found in the Alerts module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config cms-auth | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|---|----------------------|
| cms-connect- store.template - Use this template to deploy the Connect and Store module and all associate d component s. The default configuration deploys the services found in the Connect and Store module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config cms-auth cms-alerts | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|--|----------------------|
| cms-api.t emplate - Use this template to deploy the API module and all associate d component s. The default configuration deploys the services found in the API module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config cms-auth cms-connect-store | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|--|----------------------|
| cms-ev-battery-health.templ ate - Use this template to deploy the EV Battery Health module and all associate d component s. The default configuration deploys the services found in the EV Battery Health module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config cms-auth cms-api cms-alerts | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|---|----------------------|
| cms-vehicle- simulator.templ ate - Use this template to deploy the Vehicle Simulator module and all associate d component s. The default configuration deploys the services found in the Vehicle Simulator module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|---|----------------------|
| cms-fleetwise- connector.tem plate - Use this template to deploy the FleetWise Connector module and all associate d component s. The default configuration deploys the services found in the FleetWise Connector module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|--|-------------------------|---|---|----------------------|
| cms-provi sioning.t emplate - Use this template to deploy the Vehicle Provision ing module and all associate d component s. The default configuration deploys the services found in the Vehicle Provisioning module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|---|----------------------|
| cms-predictive- maintenance. template - Use this template to deploy the Predictive Maintenance module and all associate d component s. The default configuration deploys the services found in the Predictiv e Maintenan ce module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config | None |

| Template name | Download the template | Launch the template in the AWS Management Console | Prequisites | Optional Features |
|---|-------------------------|---|---|----------------------|
| cms-ui.te mplate - Use this template to deploy the Fleet Management UI module and all associate d component s. The default configuration deploys the services found in the Fleet Managemen t UI module section, and you can customize the template to meet your specific needs. | view template button | launch button | auth-setup vpc cms-config | None |
| cms-sampl e.template - Use this template for integrati ng custom modules. | view template button | launch button | None | None |



Note

AWS CloudFormation resources are created from AWS Cloud Development Kit (AWS CDK) constructs.

If you deployed a previous version of this solution, you must first uninstall your previous deployment and then deploy the new version.

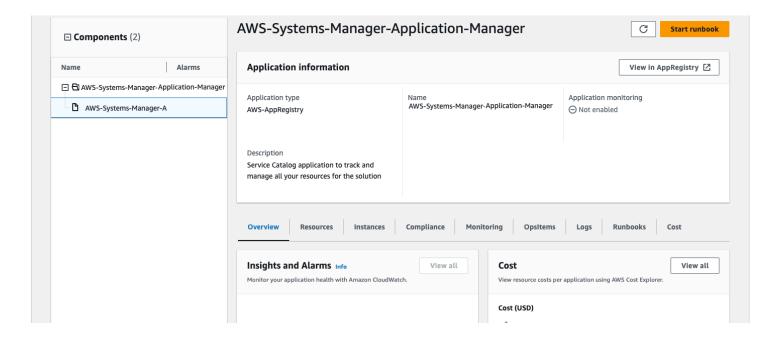
Monitor the solution with Service Catalog AppRegistry

This solution includes a Service Catalog AppRegistry resource to register the CloudFormation template and underlying resources as an application in both <u>Service Catalog AppRegistry</u> and <u>AWS</u> Systems Manager Application Manager.

AWS Systems Manager Application Manager gives you an application-level view into this solution and its resources so that you can:

- Monitor its resources, costs for the deployed resources across stacks and AWS accounts, and logs associated with this solution from a central location.
- View operations data for the resources of this solution (such as deployment status, CloudWatch alarms, resource configurations, and operational issues) in the context of an application.

The following figure depicts an example of the application view for the solution stack in Application Manager.



myApplications dashboard

Service Catalog AppRegistry integrates with the myApplications dashboard. You can access the dashboard in the **Metrics** tab in the **Catalog Item** view in Backstage. This page shows the total cost of your deployment and provides a link to the myApplications dashboard for the linked application.

myApplications dashboard 99

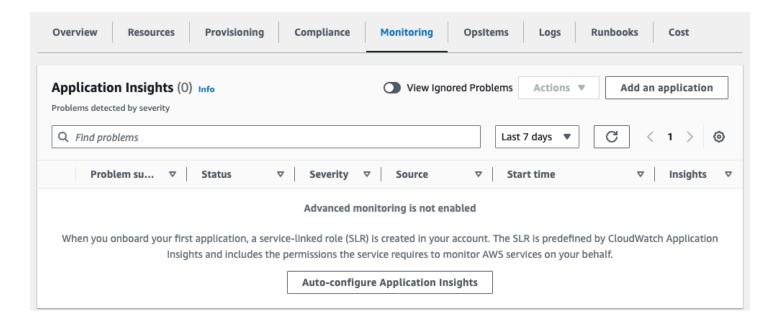
From this dashboard, you can view further cost insights and configure and view further metrics for your solution by using services such as Security Hub, CloudWatch, and Cost Explorer.

Activate CloudWatch Application Insights

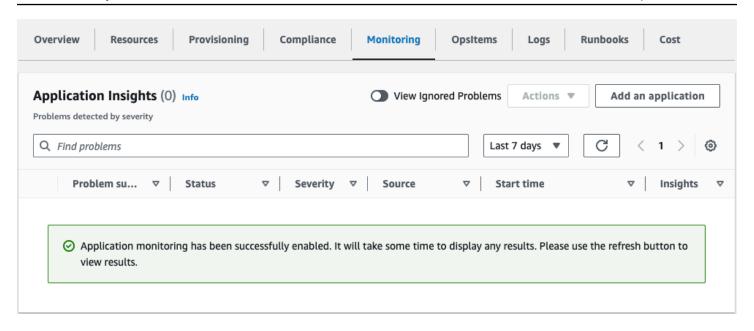
- 1. Sign in to the Systems Manager console.
- 2. In the navigation pane, choose **Application Manager**.
- 3. In **Applications**, search for the application name for this solution and select it.

The application name will have App Registry in the **Application Source** column, and will have a combination of the solution name, Region, account ID, or stack name.

- 4. In the **Components** tree, choose the application stack you want to activate.
- 5. In the Monitoring tab, in Application Insights, select Auto-configure Application Insights.



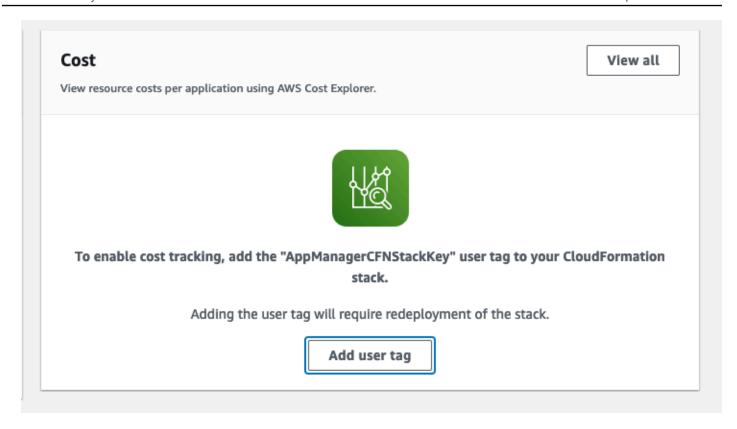
Monitoring for your applications is now activated and the following status box appears:



Confirm cost tags associated with the solution

After you activate cost allocation tags associated with the solution, you must confirm the cost allocation tags to see the costs for this solution. To confirm cost allocation tags:

- 1. Sign in to the Systems Manager console.
- 2. In the navigation pane, choose **Application Manager**.
- 3. In **Applications**, choose the application name for this solution and select it.
- 4. In the Overview tab, in Cost, select Add user tag.



5. On the Add user tag page, enter confirm, then select Add user tag.

The activation process can take up to 24 hours to complete and the tag data to appear.

Activate cost allocation tags associated with the solution

After you confirm the cost tags associated with this solution, you must activate the cost allocation tags to see the costs for this solution. The cost allocation tags can only be activated from the management account for the organization.

To activate cost allocation tags:

- 1. Sign in to the AWS Billing and Cost Management and Cost Management console.
- 2. In the navigation pane, select **Cost Allocation Tags**.
- 3. On the **Cost allocation tags** page, filter for the AppManagerCFNStackKey tag, then select the tag from the results shown.
- 4. Choose Activate.

AWS Cost Explorer

You can see the overview of the costs associated with the application and application components within the Application Manager console through integration with AWS Cost Explorer. Cost Explorer helps you manage costs by providing a view of your AWS resource costs and usage over time.

- 1. Sign in to the AWS Cost Management console.
- 2. In the navigation menu, select **Cost Explorer** to view the solution's costs and usage over time.

AWS Cost Explorer 103

Update the solution

If you deployed a previous version of this solution, you must first $\underline{\text{uninstall}}$ your previous deployment and then $\underline{\text{deploy}}$ the new version.

Troubleshooting

This section provides known issue resolution when deploying the solution. If these instructions don't address your issue, Contact AWS Support provides instructions for opening an AWS Support case for this solution.

Problem: Lambda Runtime not supported

Occasionally, an error could appear similar to Lambda Runtime (...) not supported. Supported list [...]. This is due to AWS CDK or third-party library updates that might create Lambda functions with non-supported runtime versions. To resolve this issue, you must add the necessary runtime version to the Lambda function's configuration and Lambda function aspect supported list.

Resolution

Add the necessary runtime version to the Lambda functions supported list:

- Explicitly define the runtimes for a Lambda function in its {https---docs-aws-amazon-com-cdkapi-v2-docs-aws-cdk- lib-aws-lambda-Function-html}[configuration].
- 2. Find the Lambda function aspect for the appropriate module, typically located in source/infrastructure/spects/validations.py, and add the new runtime version to the supported runtime list.

Problem: Multiple ProvisionedVehicles active certificates

Attempting to concurrently provision multiple vehicles of the same VIN can lead to a non-valid database step for the ProvisionedVehicles DynamoDB table. Specifically, the table will have multiple **ActiveCertificates** for a single VIN, meaning multiple records exist in the table. To resolve this issue, delete all records for that VIN and reprovision.

Resolution

Delete all ProvisionedVehicles table records for the invalid VIN, and reprovision the vehicle.

1. Delete all entries for the ProvisionedVehicles DynamoDB table with the specified VIN, either through the console or the AWS Command Line Interface (AWS CLI).

2. Reprovision the vehicle with one request. Reprovisioning the vehicle generates a new valid certificate, registers the vehicle with the new certificate, and creates a new record for the vehicle.

Problem: Redirect mismatch error when using Cognito Instance deployed by Auth-Setup module

When the auth-setup module is deployed using Cognito, it includes a user app client with an allowed list of Callback URLs. If not provided it takes the default value of https://example.com. The appropriate callback URLs are necessary to properly sign-in to the Backstage portal. If not configured properly, the redirect, and login attempt, will fail.

Resolution

Log in to AWS console and navigate to Cognito. Locate the cms-users app client, and assess the Callback URLs. Ensure your backstage domain is included with the pre-configured URL expected by Backstage. For more details, see the Auth Setup README.

Contact AWS Support

If you have <u>AWS Developer Support</u>, <u>AWS Business Support</u>, or <u>AWS Enterprise Support</u>, you can use the Support Center to get expert assistance with this solution. The following sections provide instructions.

Issues Or Questions:

If you encounter any issues with the code or need to search through existing or past issues, visit the <u>Issues</u> section of *Connected Mobility Solution on AWS - Issues* code repository. This section serves as a centralized platform to report new problems, track their resolution status, and review previously addressed questions and discussions.

Create case

- 1. Sign in to Support Center.
- 2. Choose Create case.

How can we help?

- 1. Choose **Technical**.
- 2. For **Service**, select **Solutions**.
- 3. For Category, select Other Solutions.
- 4. For **Severity**, select the **Severity 4** option.
- 5. When you enter the **Service**, **Category**, and **Severity**, the interface populates links to common troubleshooting questions. If you can't resolve your question with these links, choose **Next step:**Additional information.

Additional information

- 1. For **Subject**, enter text summarizing your question or issue.
- 2. For **Description**, title the description section with "Assign this ticket to CTI: AWS Solutions / Connected Mobility Solution on AWS / Support Request" and then describe the issue in detail.
- 3. Choose Attach files.
- 4. Attach the information that AWS Support needs to process the request.

Help us resolve your case faster

- 1. Enter the requested information.
- 2. Choose Next step: Solve now or contact us.

Solve now or contact us

- Review the Solve now solutions.
- 2. If you can't resolve your issue with these solutions, choose **Contact us**, enter the requested information, and choose **Submit**.

How can we help?

Uninstall the solution

This solution creates multiple CloudFormation deployments. Some resources cannot be uninstalled directly with CloudFormation and must be deleted by using the AWS Management Console or the AWS CLI.

Capture the deployment UUID

Capture and store the deployment UUIDs (Universally Unique ID) of the solution. This is used to look for any resources not destroyed by CloudFormation after teardown completes.

```
make get-acdp-deployment-uuid
make get-cms-deployment-uuid
```

The output will be uuidv4 strings, capture and store both:

XXXXXXX-XXXX-XXXX-XXXXXXXXXXXXXXX

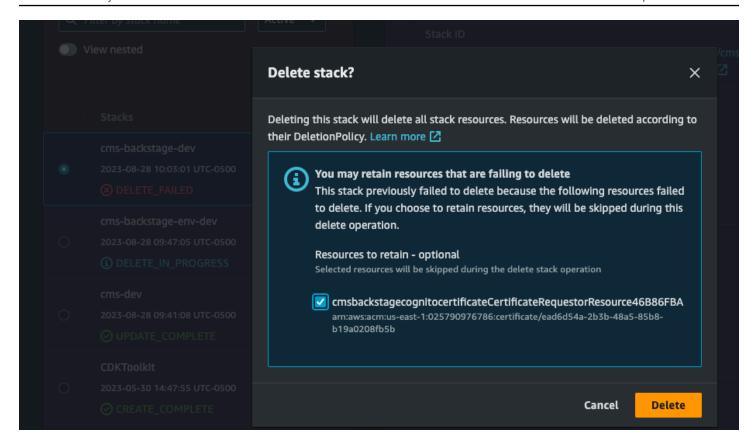
Delete CMS on AWS modules in order

make destroy



Note

Backstage might fail to delete due to the ACM certificate creation custom resource. After delete fails, select DELETE again and select retain on the custom resource. This will not leave any resources in the account.



Delete the Backstage ACM certificate (optional)

Navigate to Amazon Certificate Manager, and delete the Backstage certificate.

Manually clean up resources

You must manually clean up the following resources:

- S3 buckets
- DynamoDB tables
- Cognito user pool
- · KMS keys

Locate leftover resources using the following command, which first requires you to export the DEPLOYMENT_UUID variable using each of the values previously acquired from AWS Systems Manager.

If you tore down the ACDP stack without capturing the UUID, you can run the following command by removing the Solutions: DeploymentUUID Key filter, however the results will include other CMS on AWS stacks if they exist, so use this method with caution.

```
export DEPLOYMENT_UUID=<DEPLOYMENT_UUID_VALUE_FROM_SSM>
aws resourcegroupstaggingapi get-resources --tag-filters \
Key=Solutions:SolutionID, Values=S00241 \
Key=Solutions:DeploymentUUID, Values=$DEPLOYMENT_UUID \
--query "ResourceTagMappingList[*].ResourceARN"
```

This guery results in a list of ARNs to assist you with locating the resources in the AWS Management Console. Resources can then be manually deleted, or deleted with a script, utilizing the resource ARNs where appropriate.

Important

Some resources may take some time to clean up after CloudFormation finishes tearing down, and could show in the output even if they no longer exist. For example, Amazon VPC, Fargate, and Amazon ECS resources can remain queryable for up to 30 minutes after deletion.

Deleting the Amazon S3 buckets

This solution is configured to retain the solution-created Amazon S3 buckets if you decide to delete the AWS CloudFormation stack to prevent accidental data loss. After uninstalling the solution, you can manually delete these Amazon S3 buckets if you don't need to retain the data. Follow these steps to delete an Amazon S3 bucket.

- 1. Sign in to the Amazon S3 console.
- 2. Choose **Buckets** from the left navigation pane.
- 3. Locate the S3 buckets created by the solution.
- 4. Select an S3 bucket.
- 5. Choose **Delete**.

To delete the Amazon S3 bucket using AWS CLI, run the following command:

\$ aws s3 rb s3://<bucket-name> --force

Deleting the Amazon DynamoDB tables

This solution is configured to retain the DynamoDB tables if you decide to delete the CloudFormation stack to prevent accidental data loss. After uninstalling the solution, you can manually delete the DynamoDB tables if you do not need to retain the data. Follow these steps:

- 1. Sign in to the Amazon DynamoDB console.
- 2. Choose **Tables** from the left navigation pane.
- 3. Locate the tables created by the solution.
- 4. Select a CMS on AWS table.
- 5. Choose **Delete**.
- 6. Repeat the steps until you have deleted all of the solution's tables.

To delete the DynamoDB tables using AWS CLI, run the following command:

\$ aws dynamodb delete-table <table-name>

Deleting the Amazon CloudWatch logs

This solution retains the CloudWatch Logs if you decide to delete the CloudFormation stack to prevent against accidental data loss. After uninstalling the solution, you can manually delete the logs if you do not need to retain the data. Follow these steps to delete the CloudWatch Logs.

- 1. Sign in to the Amazon CloudWatch console.
- 2. Choose Log Groups from the left navigation pane.
- 3. Locate the log groups created by the solution.
- 4. Select one of the log groups.
- 5. Choose **Actions Delete**.

Repeat the steps until you have deleted all the solution's log groups.

To delete the CloudWatch log groups using AWS CLI, run the following command:

\$ aws logs delete-log-group --log-group-name <log-group-name>

Deleting the AWS KMS customer managed keys

This solution retains the AWS KMS customer managed keys if you decide to delete the CloudFormation stack to prevent against accidental encrypted data loss. After uninstalling the solution, you can manually delete the keys if you do not need to use them again. Follow these steps to delete the AWS KMS keys.

- 1. Sign in to the AWS KMS console.
- 2. Choose **Customer managed keys** from the left navigation pane.
- 3. Locate the keys created by the solution.
- 4. Select one of the keys.
- 5. Choose **Key actions Schedule key deletion**.
- 6. Optionally edit the Waiting period (in days) value.
- 7. Select **Confirmation**.
- 8. Choose **Schedule deletion**.

Repeat the steps until you have deleted all the solution's customer managed keys.

To delete the AWS KMS customer managed keys using AWS CLI, run the following command:

\$ aws kms schedule-key-deletion --key-id <key-id-or-arn>

Deleting the Amazon Cognito user pools

This solution retains the Amazon Cognito user pools if you decide to delete the CloudFormation stack to prevent against accidental user data loss. After uninstalling the solution, you can manually delete the user pools if you do not need to retain the users. Follow these steps to delete the user pools.

- 1. Sign in to the Amazon Cognito console.
- 2. Choose **User pools** from the left navigation pane.
- 3. Locate the user pools created by the solution.
- 4. Select one of the user pools.

- 5. Choose **Delete**.
- 6. Select **Deactivate deletion protection**.
- 7. Enter the user pool name in the second field.
- 8. Choose **Delete**.

Repeat the steps until you have deleted all the solution's user pools.

To delete the Amazon Cognito user pool using AWS CLI, run the following command:

```
$ aws cognito-idp delete-user-pool --user-pool-id <user-pool-id>
```

Deleting the Amazon Relational Database Service snapshots

This solution retains the <u>Amazon Relational Database Service</u> (Amazon RDS) snapshots if you decide to delete the AWS CloudFormation stack to prevent against accidental data loss. After uninstalling the solution, you can manually delete the snapshots if you do not need to retain the data. Follow these steps to delete the snapshots.

- 1. Sign in to the Amazon RDS console.
- 2. Choose **Snapshots** from the left navigation pane.
- 3. Locate the snapshots created by the solution.
- 4. Select one of the snapshots.
- 5. Choose **Actions Delete snapshot**.
- 6. Choose Delete.

Repeat the steps until you have deleted all the solution's snapshots.

To delete the Amazon RDS snapshot using AWS CLI, run the following command:

```
$ aws rds delete-db-snapshot --db-snapshot-identifier <db-snapshot-identifier>
```

Developer guide

This section provides the source code for the solution and additional customizations.

Source code

Visit our <u>GitHub repository</u> to download the source files for this solution and to share your customizations with others.

The <u>AWS Cloud Development Kit (AWS CDK) (AWS CDK)</u> generates the Connected Mobility Solution on AWS templates. See the <u>README.md</u> file for additional information.

Integrating custom modules

We have provided a CMS Sample module which serves as a template for existing and future CMS on AWS modules. See the code and README.md in the <u>GitHub repository</u> for more information on development.

For more information on integrating a CMS on AWS module with the Backstage deployment, see the Backstage documentation.

Global system integrator (GSI) and independent software (ISV) integration

Global system integrators (GSIs) and independent software vendors (ISVs) can leverage CMS to build and deploy mobility services quickly and efficiently. CMS provides a modular, scalable architecture that integrates various AWS services for fleet management, telematics, and vehicle data collection. The following are key ways that GSIs and ISVs can use, customize, and white-label the solution.

Using CMS as a foundation

CMS is designed to serve as a flexible foundation for connected mobility solutions. GSIs and ISVs can use the core functionality provided by CMS to build services for managing fleets, monitoring vehicle health, and processing telematics data. The solution provides:

Source code 114

- **Vehicle connectivity** The solution supports data ingestion from both internal combustion and electric vehicles that use IoT technologies.
- **Data processing and storage** The solution uses AWS IoT services, Lambda, and DynamoDB to process and store data efficiently. This provides a scalable platform for future expansion.
- **Visualization and insights** CMS includes built-in dashboards that you can customize to provide real-time vehicle insights, metrics, and performance reports.

GSIs and ISVs can extend these features based on their specific customer requirements without building from scratch. This helps to ensure rapid time-to-market for new mobility services.

Integrating with existing systems

We built CMS with extensibility in mind, helping GSIs and ISVs to integrate the solution into existing IT and mobility infrastructures. Using a microservices-based architecture, you can connect CMS to other enterprise systems, such as:

- Enterprise resource planning (ERP) systems For managing fleet and maintenance scheduling.
- **Customer relationship management (CRM) systems** For tracking customer and vehicle interactions.
- Telematics providers –To ingest and process additional data streams.

You can facilitate integration by using the APIs provided by the AWS services within CMS, such as <u>AWS IoT Core</u> and <u>Lambda</u>. Additionally, you can enhance the solution by adding custom AWS services to address unique business needs.

White-labeling CMS for end customers

GSIs and ISVs can fully white-label CMS to meet customer branding requirements. By white-labeling, you can deliver a custom-branded mobility service solution that aligns with customer expectations, without altering the core CMS architecture. This involves:

- Customizing user interfaces The CMS frontend UI components are hosted on Amplify and Amazon S3. You can customize the frontend UI to reflect customer branding guidelines. You can also modify color schemes, logos, and styles.
- **Custom domains** Deploy the solution on customer-owned domains by using Route 53 and SSL certificates to align the white-label solution with your customer's brand identity.

• **Tailoring reports and dashboards** – Customize data visualization in the dashboards to reflect customer-specific KPIs and reporting requirements, providing insights relevant to the end users.

Extending functionality with custom modules

GSIs and ISVs can extend CMS functionality by <u>building additional modules</u>. The microservice architecture supports the development of new services using Lambda, API Gateway, and other serverless technologies. Some potential extensions include:

- Advanced vehicle health monitoring Create additional monitoring modules that use custom algorithms for predictive maintenance and diagnostics.
- **Telematics data analytics** Implement additional data analytics pipelines for advanced insights into vehicle and fleet performance.
- Integration with third-party services Build connectors to third-party services such as insurance providers, leasing companies, or government regulatory bodies.

Developing new modules follows standard AWS development practices. You can implement new modules by using the existing codebase on GitHub as a reference.

Managing deployment and operations

CMS leverages AWS cloud-native services for easy deployment, scaling, and management. GSIs and ISVs can deploy CMS for their customers by using CloudFormation templates available in the <u>GitHub repository</u>. This helps you to rapidly provision and configure the required AWS resources. You can manage ongoing operations such as scaling, monitoring, and logging through:

- **CloudWatch** For monitoring solution health and setting up alarms.
- Lambda To scale serverless services based on-demand without manual intervention.
- **DynamoDB** To automatically scale the data layer to meet storage requirements.

By automating operational tasks, GSIs and ISVs can reduce operational overhead while ensuring high availability and performance.

Collaborating with AWS for customizations

GSIs and ISVs interested in extensive customizations or bespoke implementations can work closely with AWS teams for architecture reviews, scalability assessments, and custom development

support. The flexibility of CMS helps you to introduce new capabilities in line with customer requirements. This can result in a tailored solution that scales alongside your customer's business.

For more technical details, refer to the source code on <u>GitHub</u>, which provides guidance on customization, deployment, and integration patterns.

Reference

This section includes a list of builders and managers who contributed to this solution.

Anonymized data collection

This solution includes an option to send anonymized operational metrics to AWS. We use this data to better understand how customers use this solution and related services and products. When activated, the following information is collected and sent to AWS:

- Solution ID The AWS solution identifier
- Unique ID (UUID) Randomly generated, unique identifier for each deployment
- Resource Utilization Total storage utilized by the solution
- · API Usage Usage amount of API calls to the solution's public facing API

AWS owns the data gathered through this survey. Data collection is subject to the <u>Privacy Notice</u>. To opt out of this feature, complete the following steps before launching the AWS CloudFormation template.

- 1. In the <u>source/modules/cms_config/source/infrastructure/constructs/metrics.py</u> file, replace the send_anonymous_usage_condition value with No.
- 2. Continue with the deployment as specified in the Deploy the solution section.

Contributors

- Alex Sansone
- Anthony McIntosh
- Guru Kumar
- Harshit Patel
- Kevin Hargita
- Matt Wise
- Narmeen Ali
- Saif Shaikh

Anonymized data collection 118

Revisions

Check the <u>CHANGELOG.md</u> file in the GitHub repository to see all notable changes and updates to the software. The changelog provides a clear record of improvements and fixes for each version.

Notices

Customers are responsible for making their own independent assessment of the information in this document. This document: (a) is for informational purposes only, (b) represents AWS current product offerings and practices, which are subject to change without notice, and (c) does not create any commitments or assurances from AWS and its affiliates, suppliers, or licensors. AWS products or services are provided "as is" without warranties, representations, or conditions of any kind, whether express or implied. AWS responsibilities and liabilities to its customers are controlled by AWS agreements, and this document is not part of, nor does it modify, any agreement between AWS and its customers.

Connected Mobility Solution on AWS is licensed under the terms of the Apache License Version 2.0.